

Cache Efficient Parallel Partition Algorithms

An In-Place Exclusive Read/Write Memory Algorithm

OUR RESEARCH QUESTION

Can we create an algorithm with *theoretical guarantees* that is *fast in practice*?


RESULT

We created the *Smoothed Striding Algorithm*.

Key Features:

- ▶ linear work and polylogarithmic span
(like the Standard Algorithm)
- ▶ fast in practice
(like the Strided Algorithm)
- ▶ theoretically optimal cache behavior
(unlike any past algorithm)

SMOOTHED STRIDING ALGORITHM’S PERFORMANCE



STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]


- ▶ Good cache behavior in practice
- ▶ Worst case span is $T_\infty \approx n$
- ▶ On random inputs span is $T_\infty = \tilde{O}(n^{2/3})$

Smoothed-Striding Algorithm

- ▶ Provably optimal cache behavior
- ▶ Span is $T_\infty = O(\log n \log \log n)$ with high probability in n
- ▶ Uses randomization *inside* the algorithm


SMOOTHED STRIDING ALGORITHM

Logically partition the array into chunks of adjacent elements.




Form groups U_i that contain a random element from each chunk.

This randomization step was one of our key insights; it guarantees that the U_i 's have similar compositions regardless of the input.




Perform serial partitions on each U_i in parallel over the U_i 's.


This step is highly parallel.



Define v_i = index of first element greater than the pivot in U_i .




Identify leftmost and rightmost v_i . Note that $A[k] \leq \text{pivot}$ for all $k < v_{\min}$, and $A[k] > \text{pivot}$ for all $k \geq v_{\max}$.



Recursively partition the subarray.

This step was previously impossible; adding randomization enables this step, which enables our algorithm's low span.




A KEY CHALLENGE


How do we store the U_i 's if they are all random?

Storing which elements make up each U_i takes too much space!

Strided Algorithm P_i .



Smoothed-Striding Algorithm U_i .



HOW TO STORE THE GROUPS

Key Insight: While each U_i does need to contain a random element from each chunk, the U_i 's don't need to be *independent*.

We store U_1 , and all other groups are determined by a "circular shift" of U_1 (wraparound within each chunk).

