

Cache Efficient Parallel Partition Algorithms

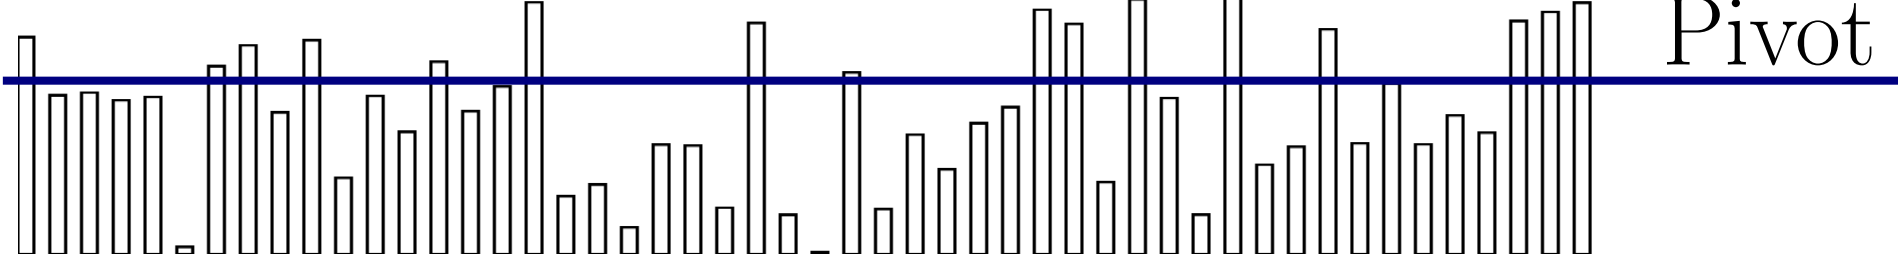
An In-Place Exclusive Read/Write Memory Algorithm

WHAT IS THE PARTITION PROBLEM?

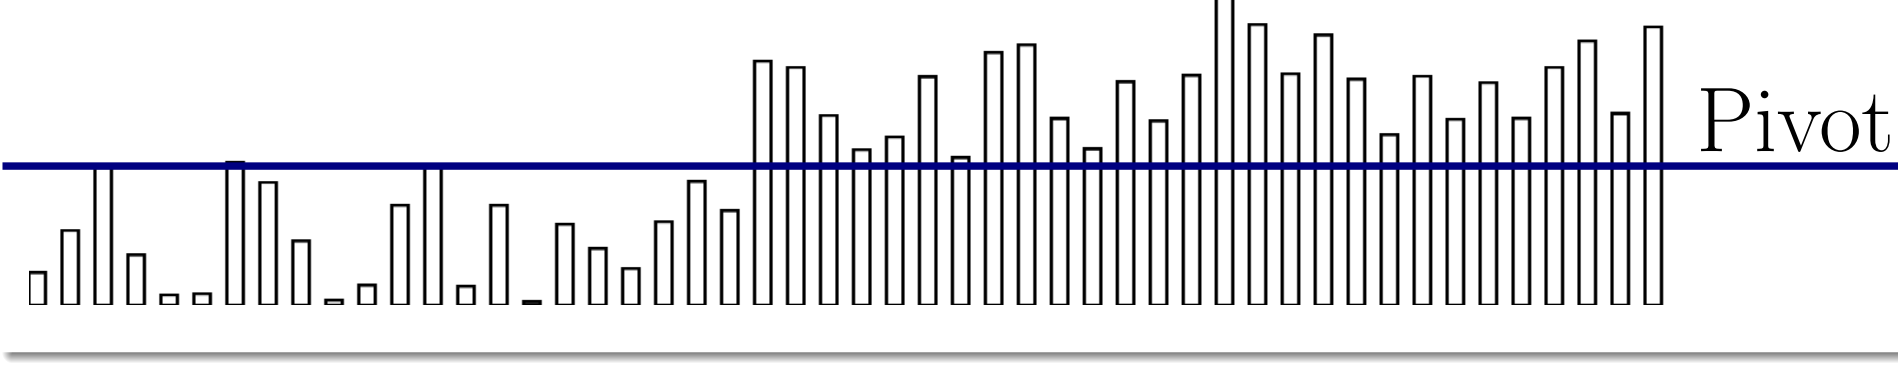
Explanation: The *Partition Problem* is to reorder the elements in a list so that elements in the same group occur in the same part of the list.

Example: A common way of grouping elements is based on whether they exceed or fall short of a certain “pivot” value.

An unpartitioned array:



An array partitioned relative to the pivot value:



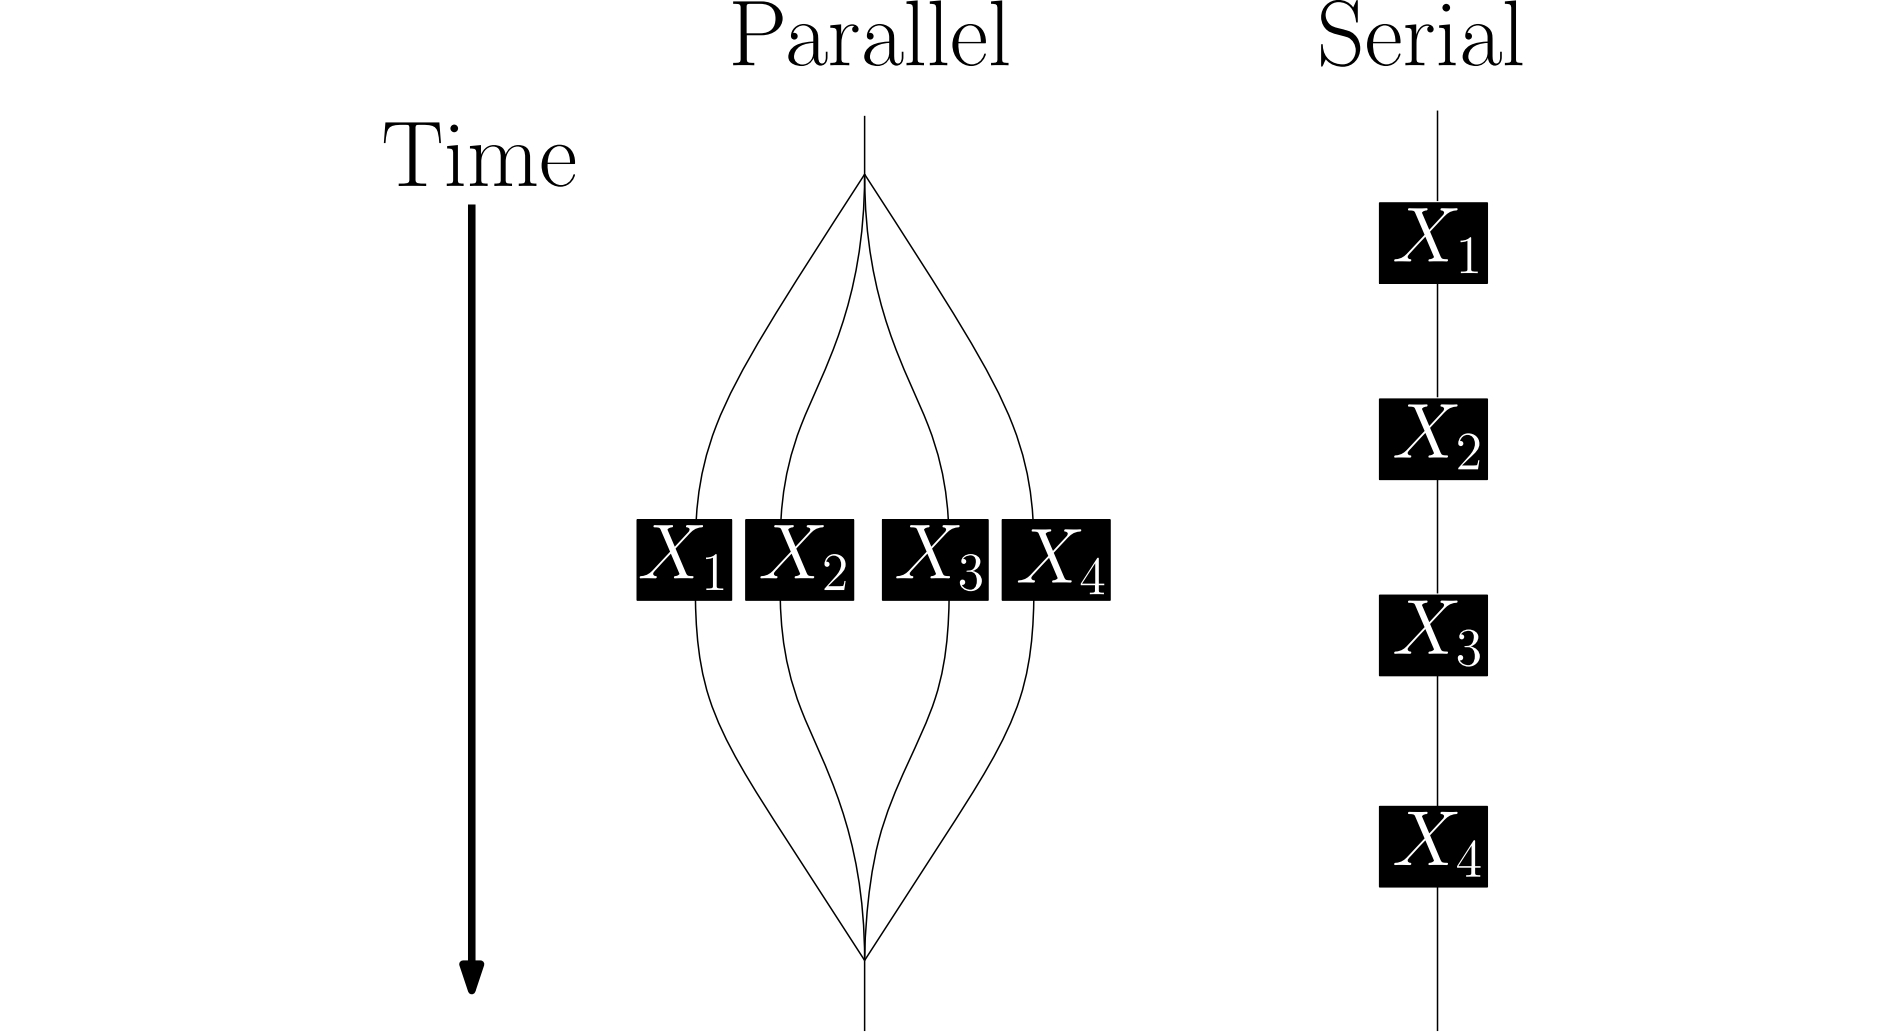
WHAT IS A PARALLEL ALGORITHM?

Explanation: Whereas a typical (i.e. serial) algorithm runs on a single processor, a *parallel algorithm* runs on $p \geq 1$ processors.

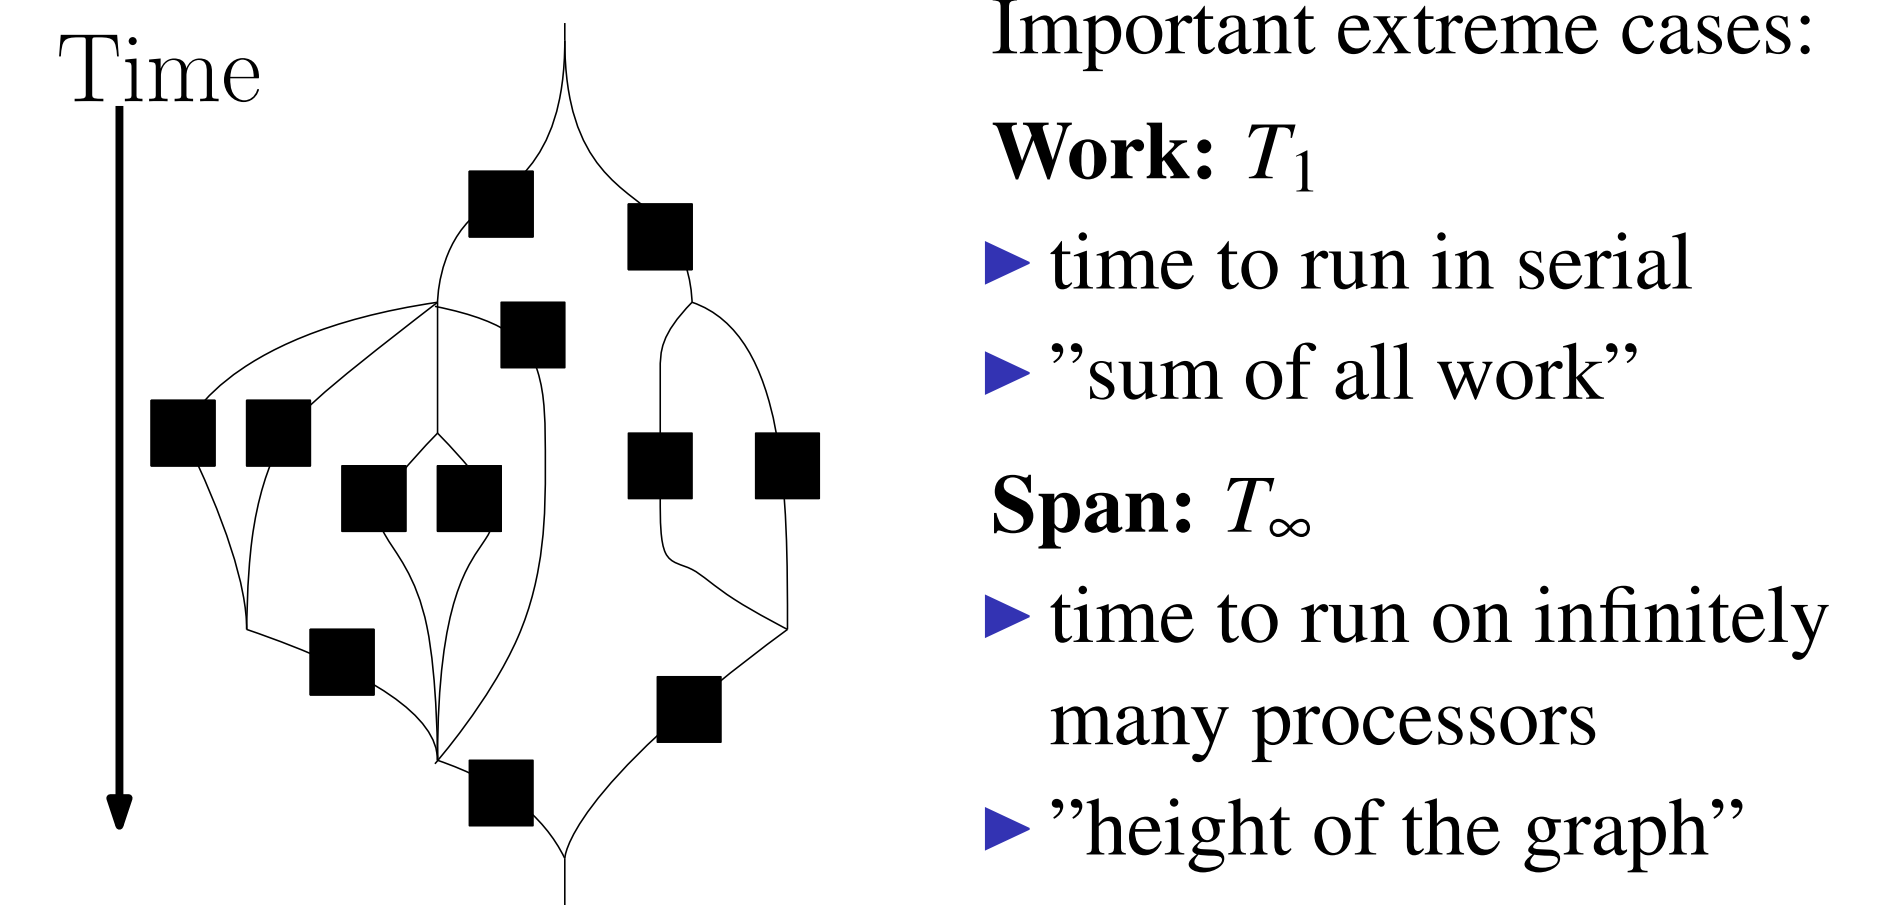
In our model of parallelism, we only allow the concurrency mechanism of parallel-for-loops; in particular our algorithm doesn’t make concurrent writes to data (it is **EREW**): we don’t allow locks or atomic variables. Being EREW is desirable because theoretical predictions apply more readily to them, and because EREW algorithms are hardware independent.

Example: Many tasks have parts that can be performed concurrently; such tasks can be performed faster with parallel computing.

Program execution in serial and in parallel:



PERFORMANCE METRICS FOR PARALLEL ALGORITHMS



Important extreme cases:

- Work:** T_1
 - ▶ time to run in serial
 - ▶ “sum of all work”
- Span:** T_∞
 - ▶ time to run on infinitely many processors
 - ▶ “height of the graph”

WHAT IS CACHE EFFICIENCY?

Explanation: *Cache* is a small part of memory that can be accessed much faster than ordinary RAM. When data is already loaded into Cache a program can rapidly access it; this is called a *cache hit*. When data needed by a program isn’t in cache it must be loaded into cache; this is called a *cache miss*, and takes time.

Remark: An algorithm with very few cache misses is *Cache Efficient*; cache efficiency leads to faster performance in practice.

Factors Affecting Cache-Efficiency:

- ▶ Perform low number of passes over the data
- ▶ Don’t use extra memory, i.e. are *In-Place*
- ▶ Deal with elements that are close in memory together

PREVIOUS WORK ON THE PARTITION PROBLEM

The “Standard Algorithm” is **theoretically optimal with span $O(\log n)$, but slow in practice due to poor cache behavior.**

The **fastest algorithms in practice lack theoretical guarantees**

▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

Not Exclusive Read/Write Memory

▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

No locks or atomic-variables, but no bound on span

WHY IS THE PARTITION PROBLEM IMPORTANT?

The Partition Problem is a fundamental problem in computer science. Additionally, it is a subproblem that must be solved in many algorithms such as:

▶ *Parallel Quicksort* (the most prominent application of partition algorithms)

▶ *Filtering operations*

Furthermore, the partition problem is of great practical importance as:

▶ Humans want organized data often, e.g. performing “ORDER BY” on information from a database, or simply ordering data in a spreadsheet

▶ Many algorithms run faster, or rely on, sorted data

OUR RESEARCH QUESTION

Can we create an algorithm with *theoretical guarantees* that is *fast in practice*?

RESULT

We created the *Smoothed Striding Algorithm*.

Key Features:

▶ linear work and polylogarithmic span

(like the Standard Algorithm)

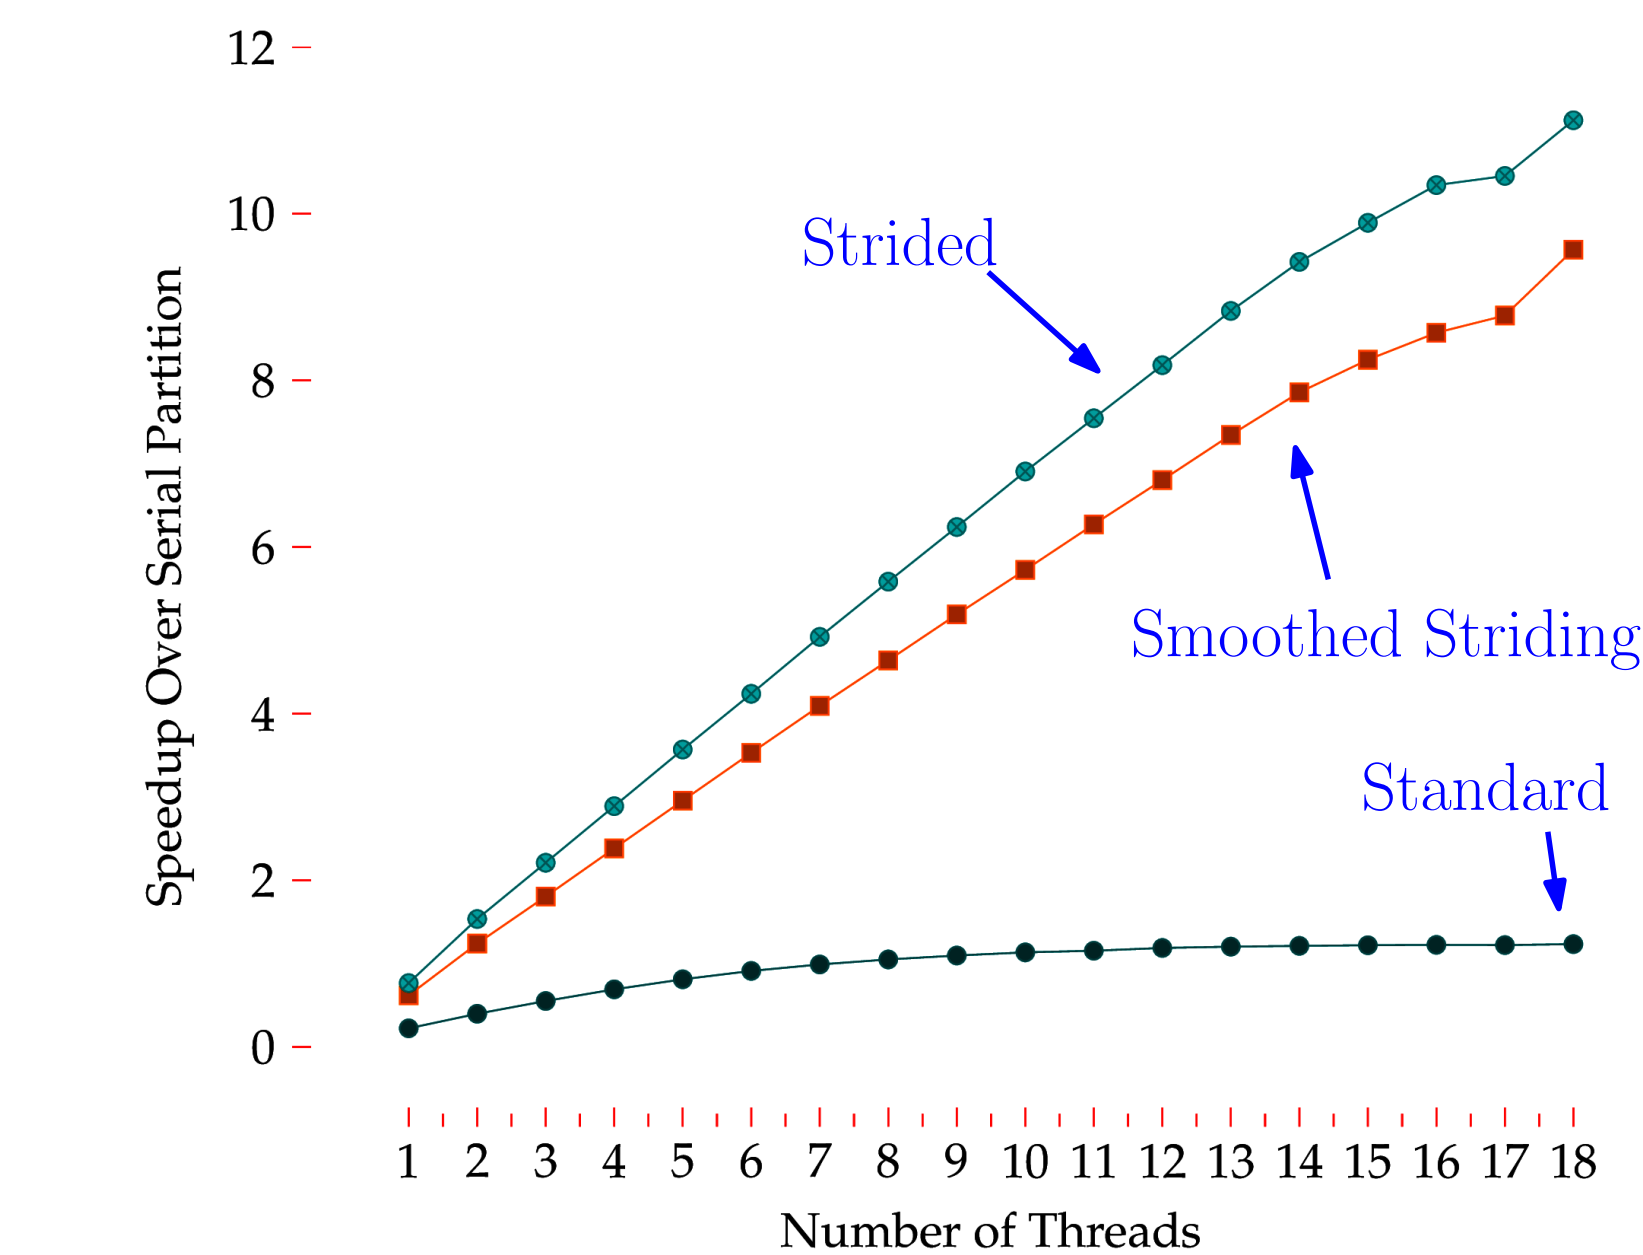
▶ fast in practice

(like the Strided Algorithm)

▶ theoretically optimal cache behavior

(unlike any past algorithm)

SMOOTHED STRIDING ALGORITHM’S PERFORMANCE



STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

▶ Good cache behavior in practice

▶ Worst case span is $T_\infty \approx n$

▶ On random inputs span is $T_\infty = \tilde{O}(n^{2/3})$

Smoothed-Striding Algorithm

▶ Provably optimal cache behavior

▶ Span is $T_\infty = O(\log n \log \log n)$ with high probability in n

▶ Uses randomization *inside* the algorithm

APPLICATION TO PARALLEL QUICKSORT

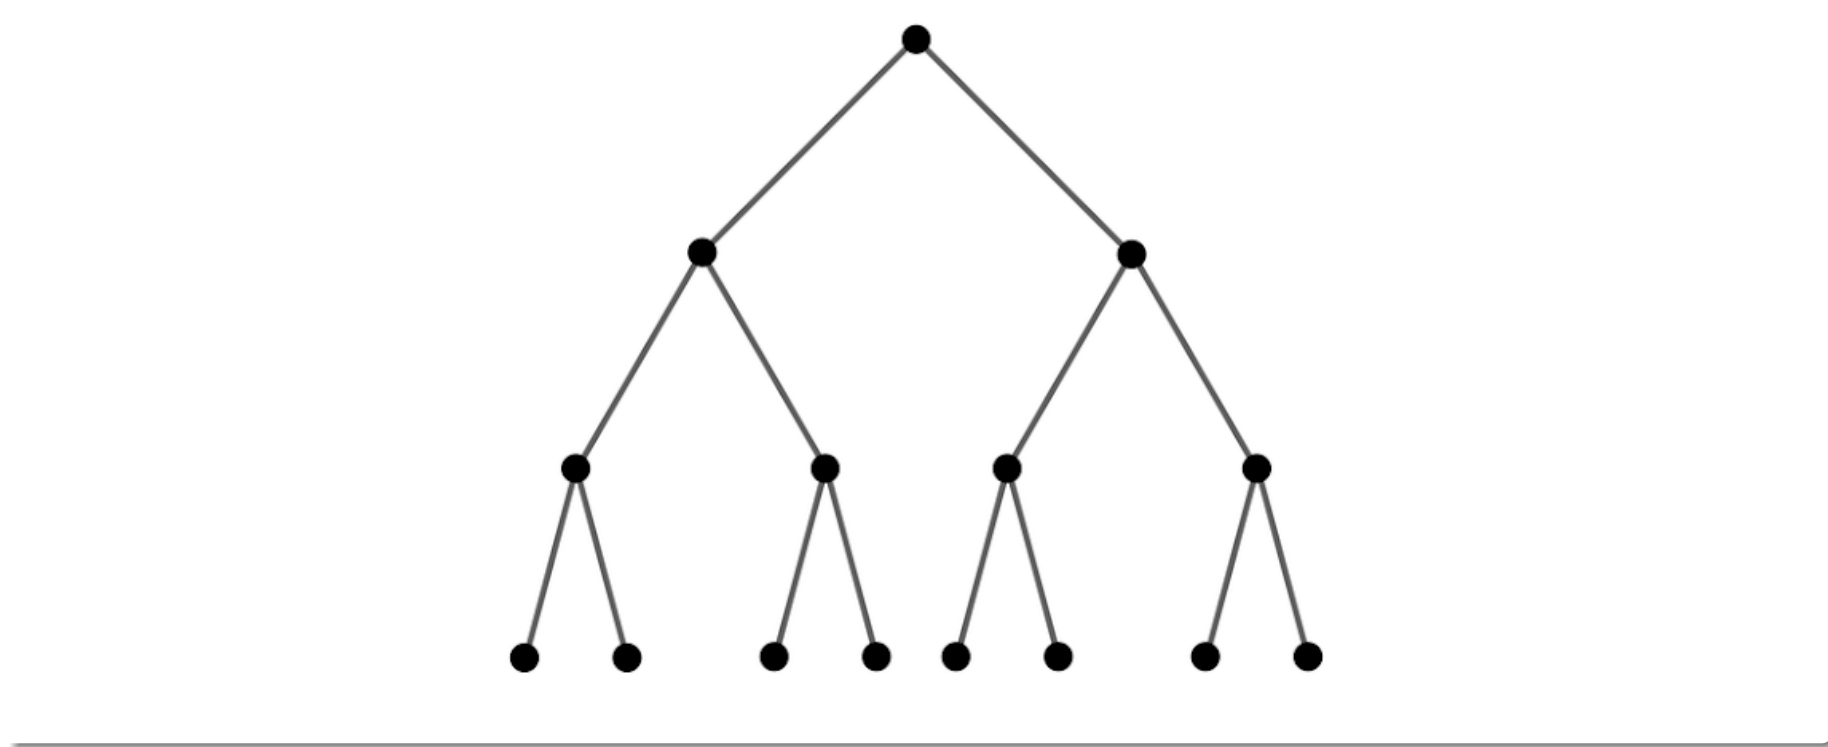
Parallel Quicksort is the most important application of Parallel Partition. Parallel Quicksort works as follows:

▶ Chose a pivot value randomly from the array

▶ *Partition* the array relative to the pivot value

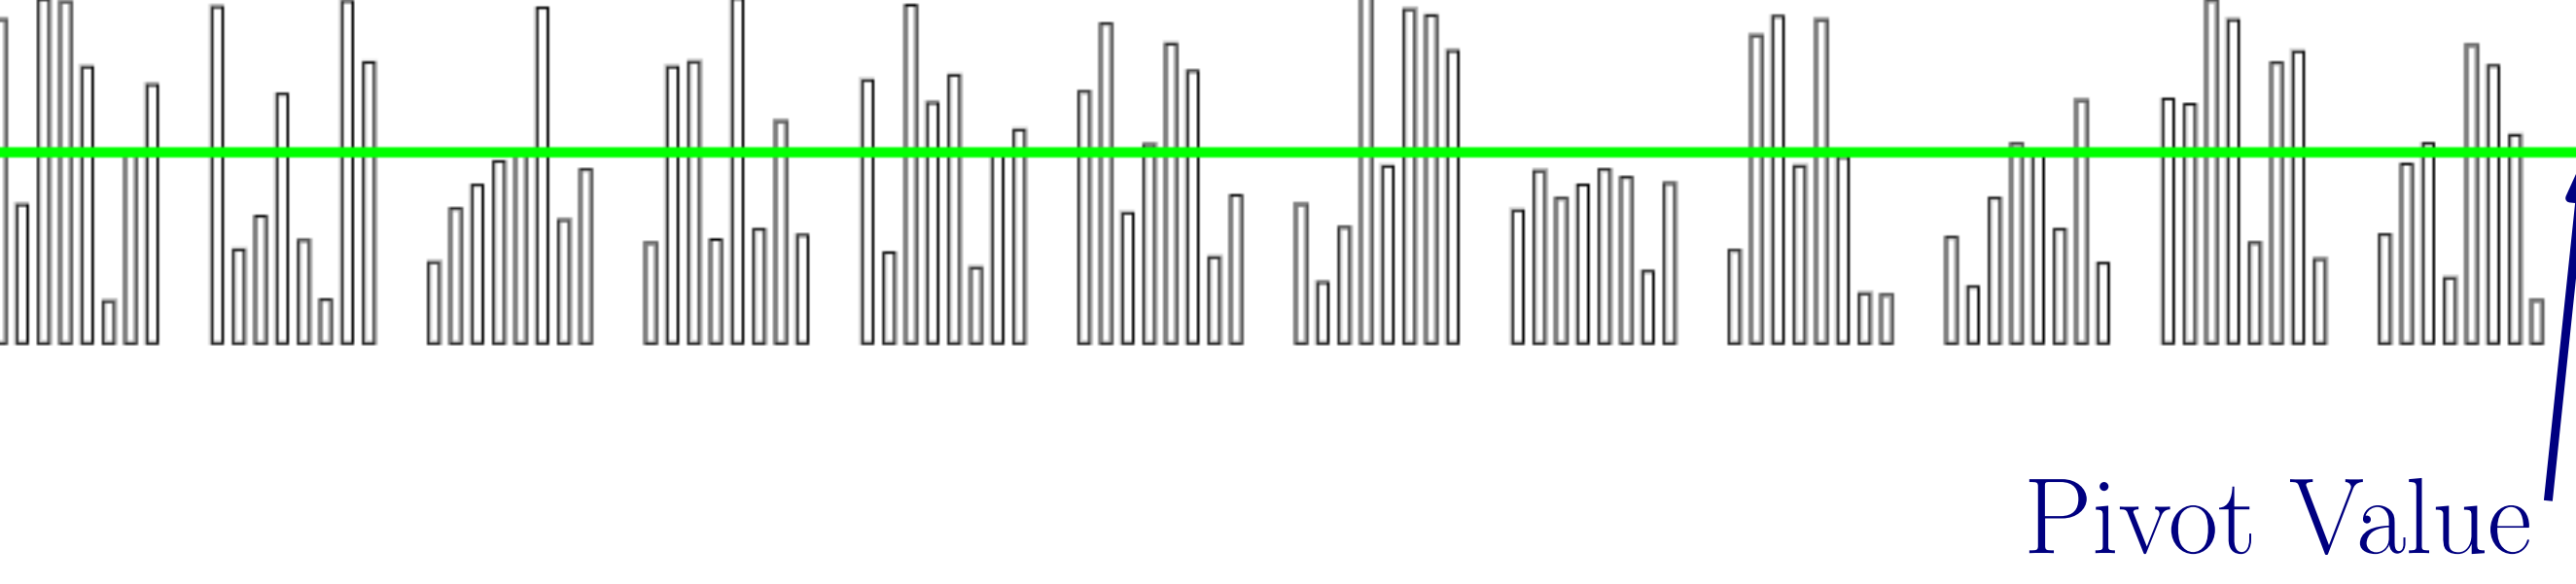
▶ Recursively sort the subarrays (in parallel)

The depth of recursion is $O(\log n)$ with high probability in n , and each level of recursion requires span $O(\log n \log \log n)$ when using the Smoothed Striding algorithm. This results in span $O(\log^2 n \log \log n)$ and work $O(n \log n)$ for the entire Parallel Quicksort – which is within a factor of $\log \log n$ of optimal span – while additionally guaranteeing cache-friendliness.



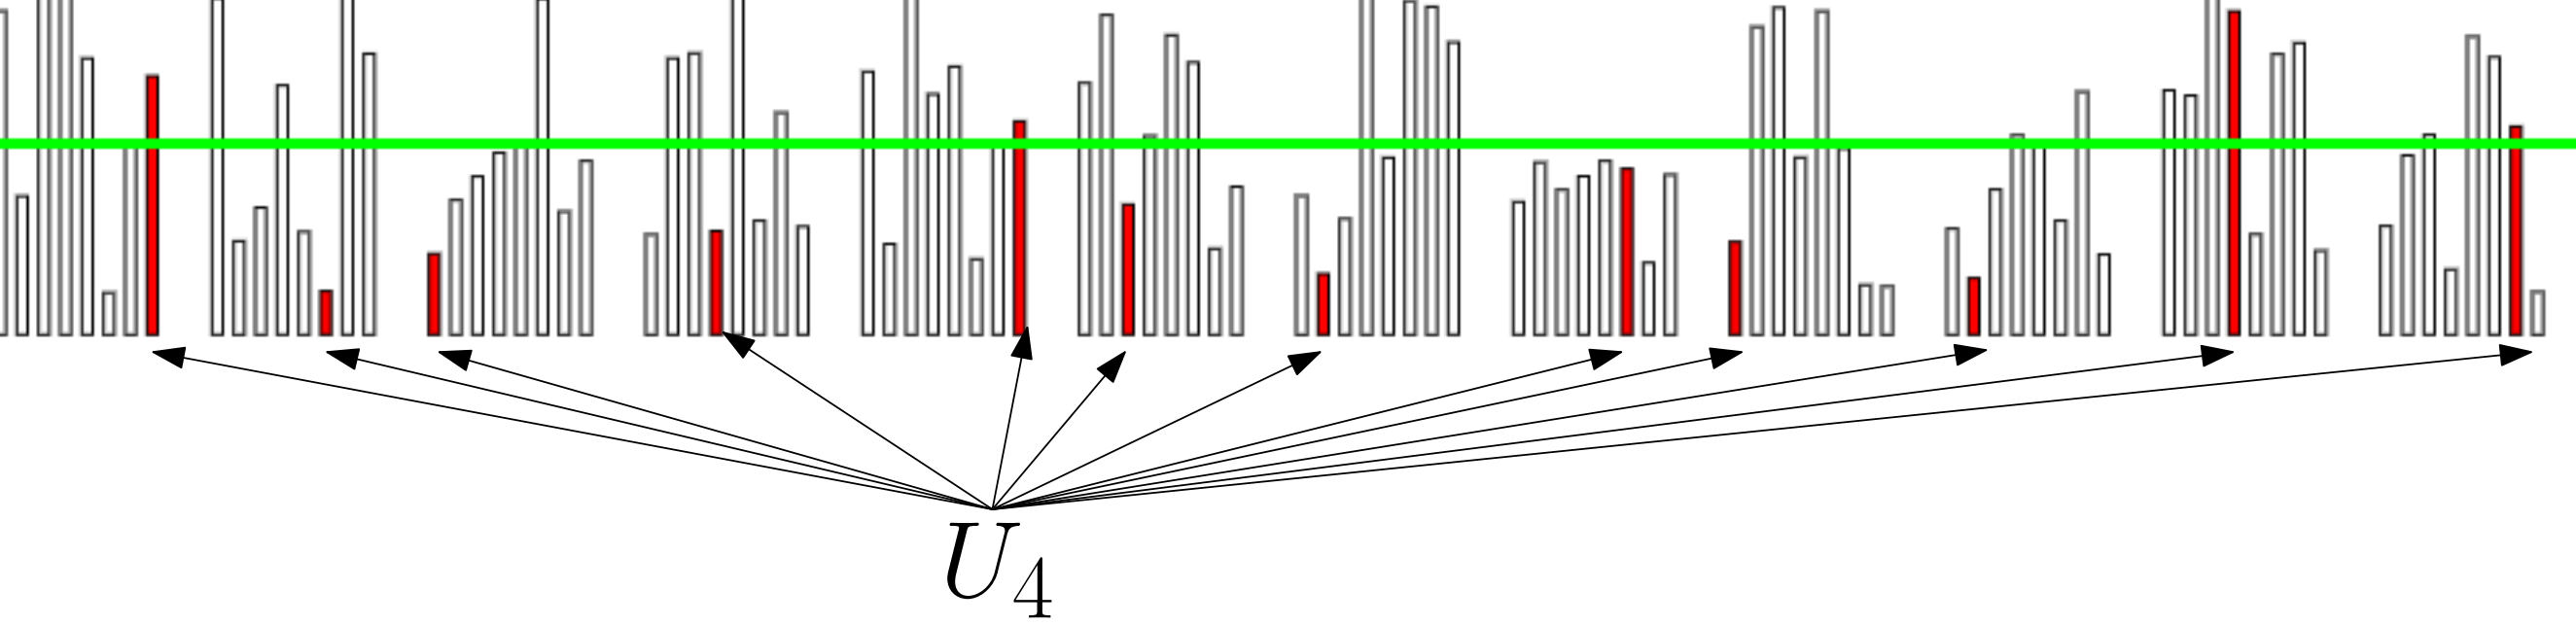
SMOOTHED STRIDING ALGORITHM

Logically partition the array into chunks of adjacent elements.



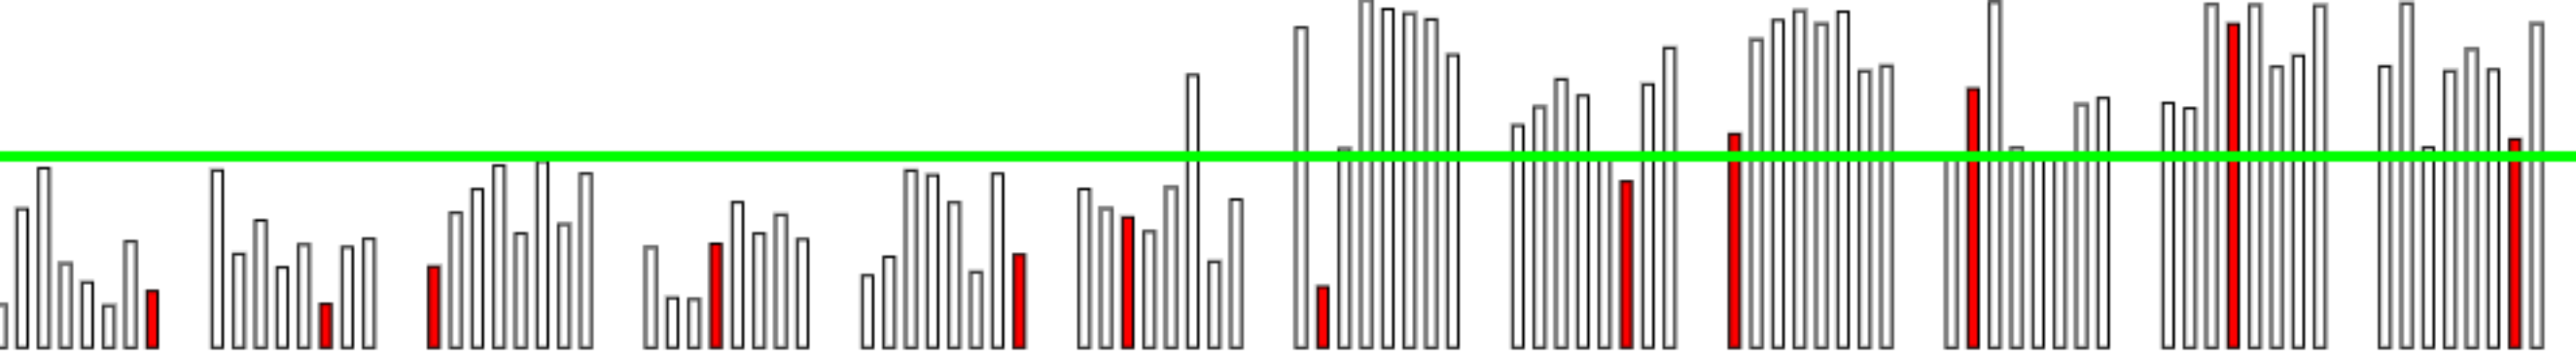
Form groups U_i that contain a random element from each chunk.

This randomization step was one of our key insights; it guarantees that the U_i 's have similar compositions regardless of the input.

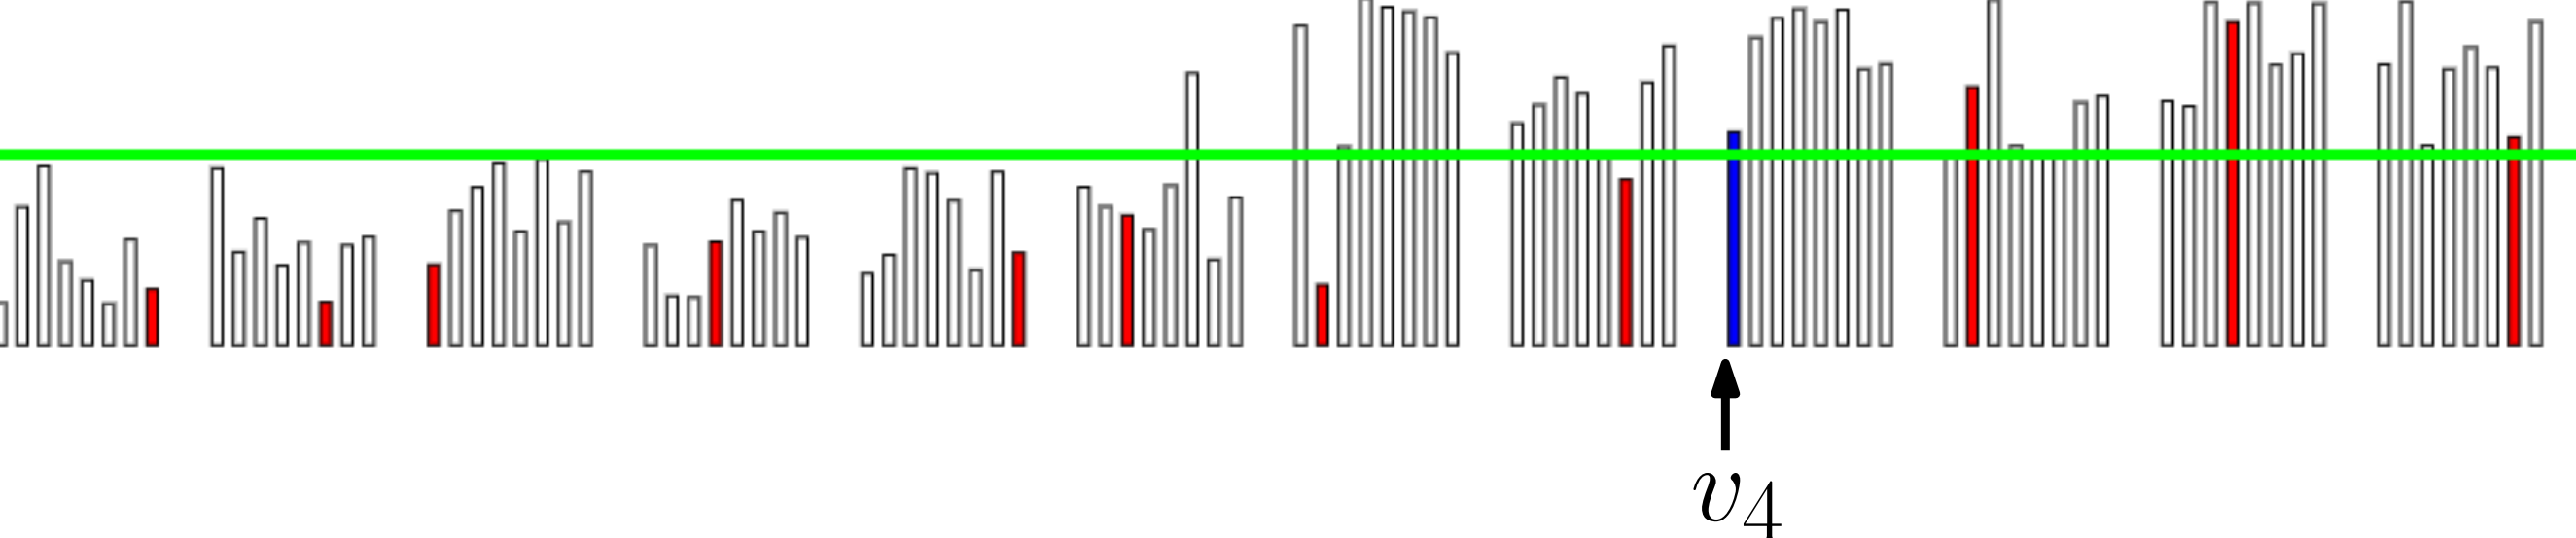


Perform serial partitions on each U_i in parallel over the U_i 's.

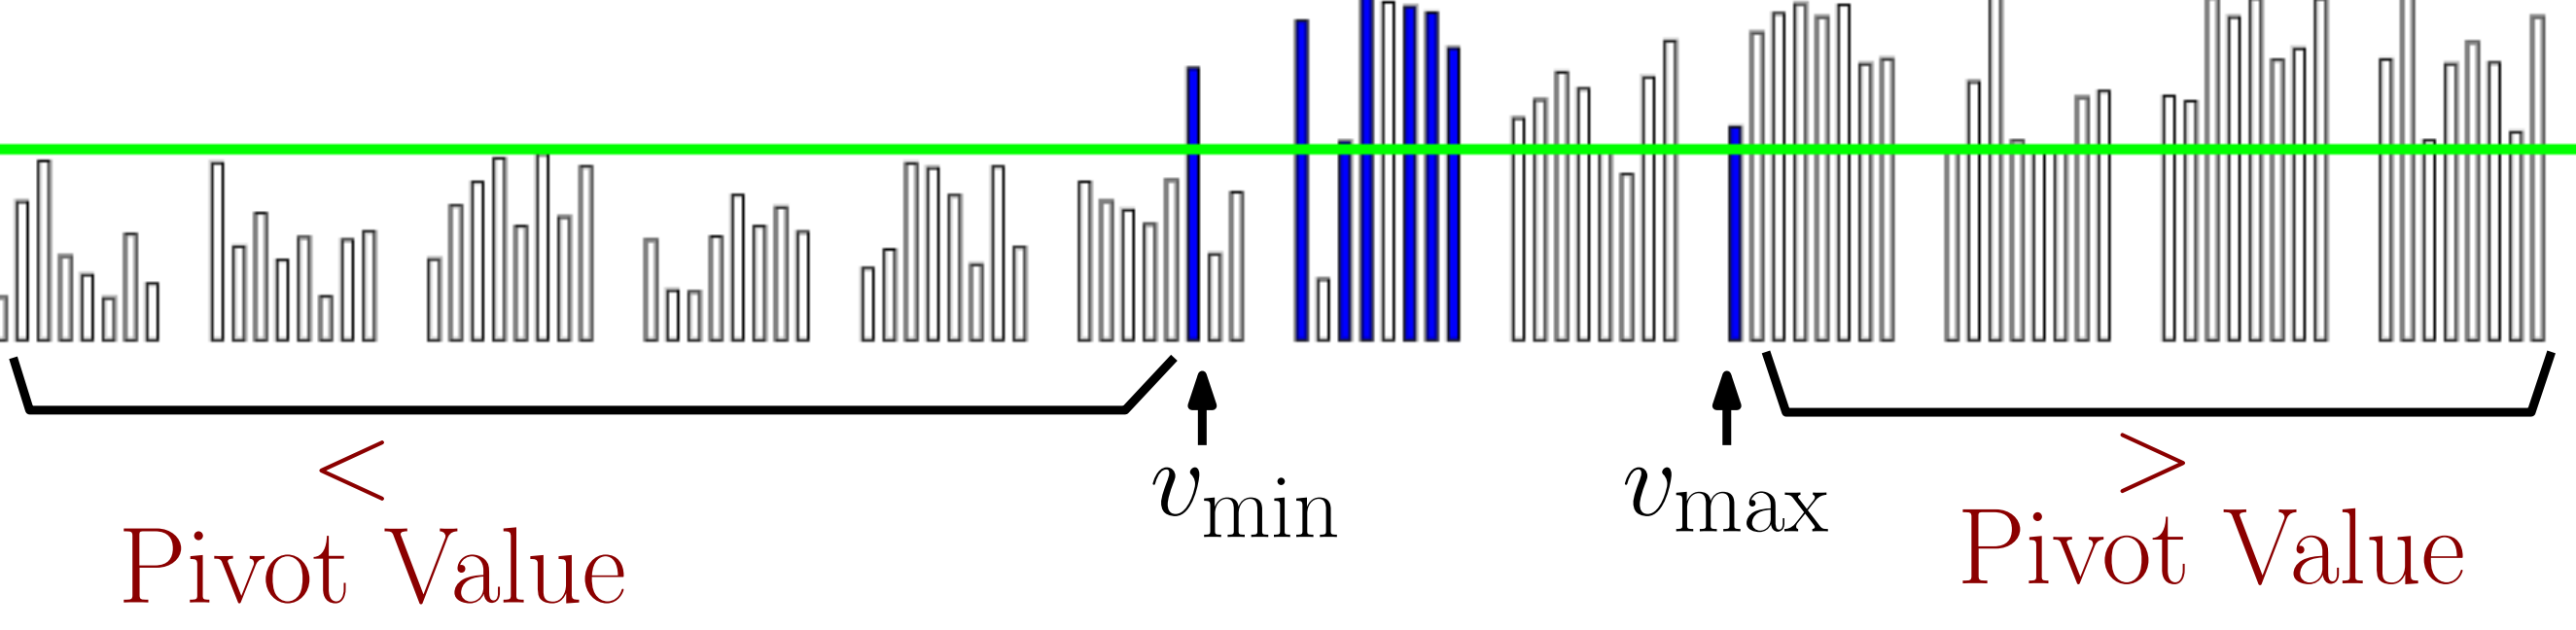
This step is highly parallel.



Define v_i = index of first element greater than the pivot in U_i .

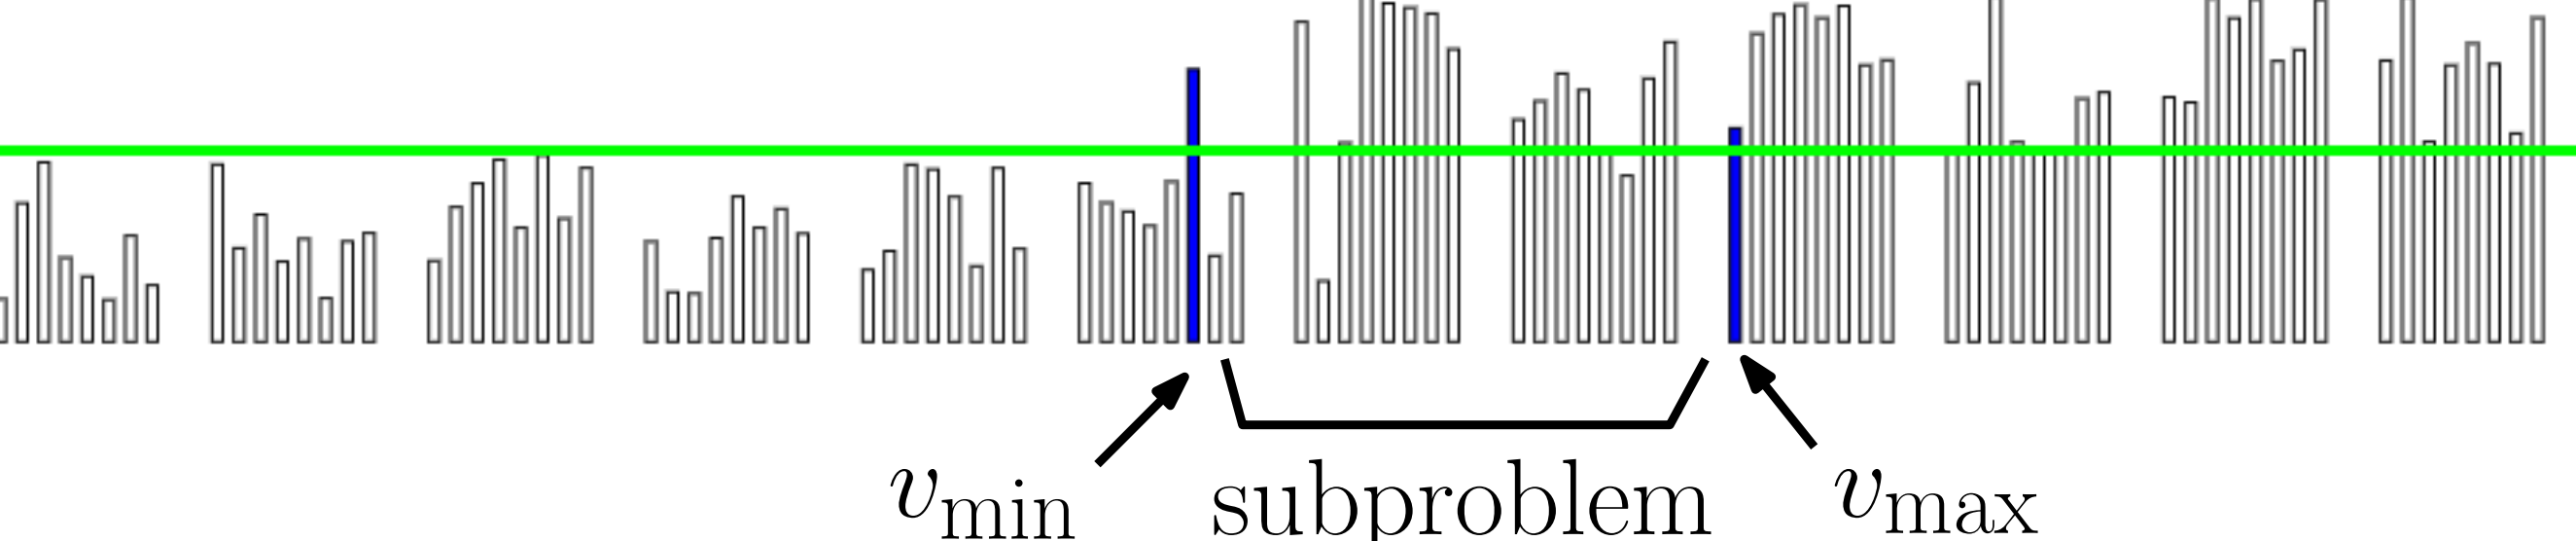


Identify leftmost and rightmost v_i . Note that $A[k] \leq \text{pivot}$ for all $k < v_{\min}$, and $A[k] > \text{pivot}$ for all $k \geq v_{\max}$.



Recursively partition the subarray.

This step was previously impossible; adding randomization enables this step, which enables our algorithm’s low span.

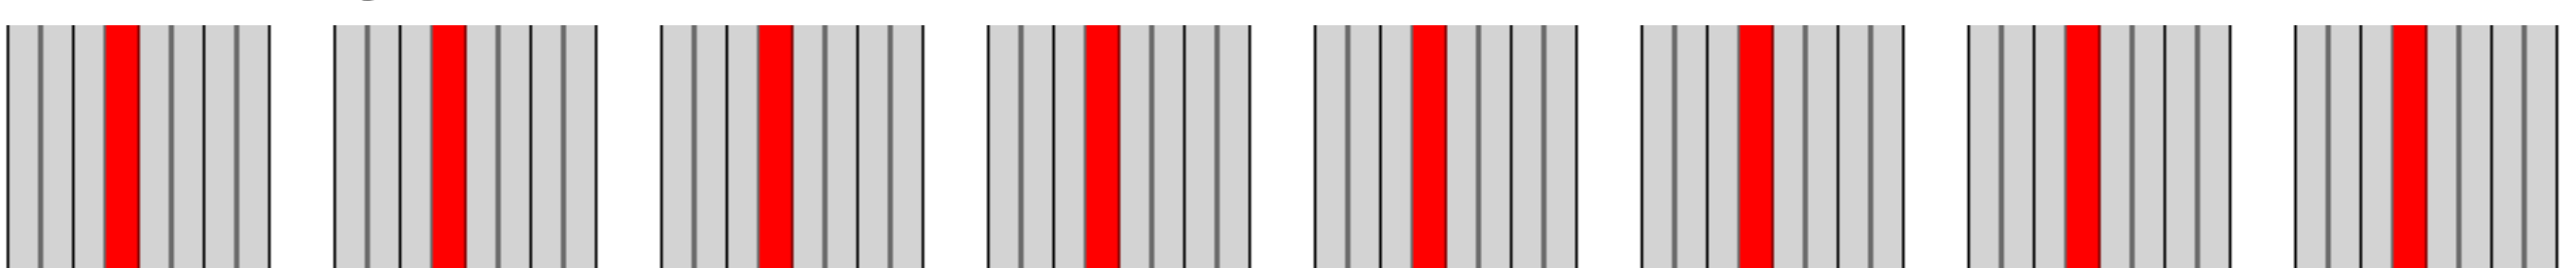


A KEY CHALLENGE

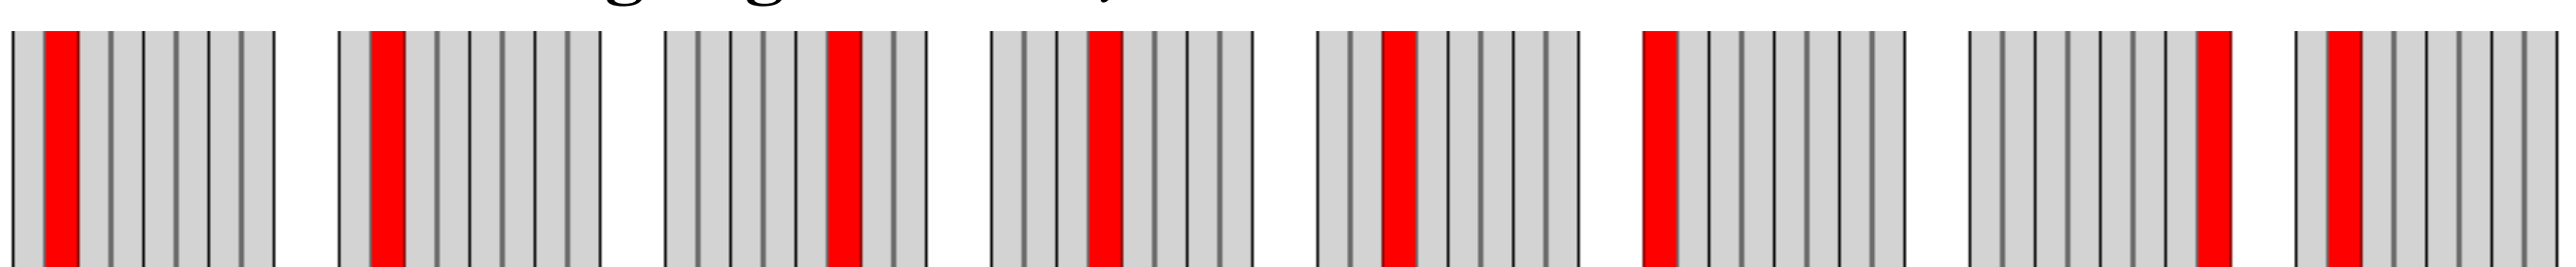
How do we store the U_i 's if they are all random?

Storing which elements make up each U_i takes too much space!

Strided Algorithm P_i .



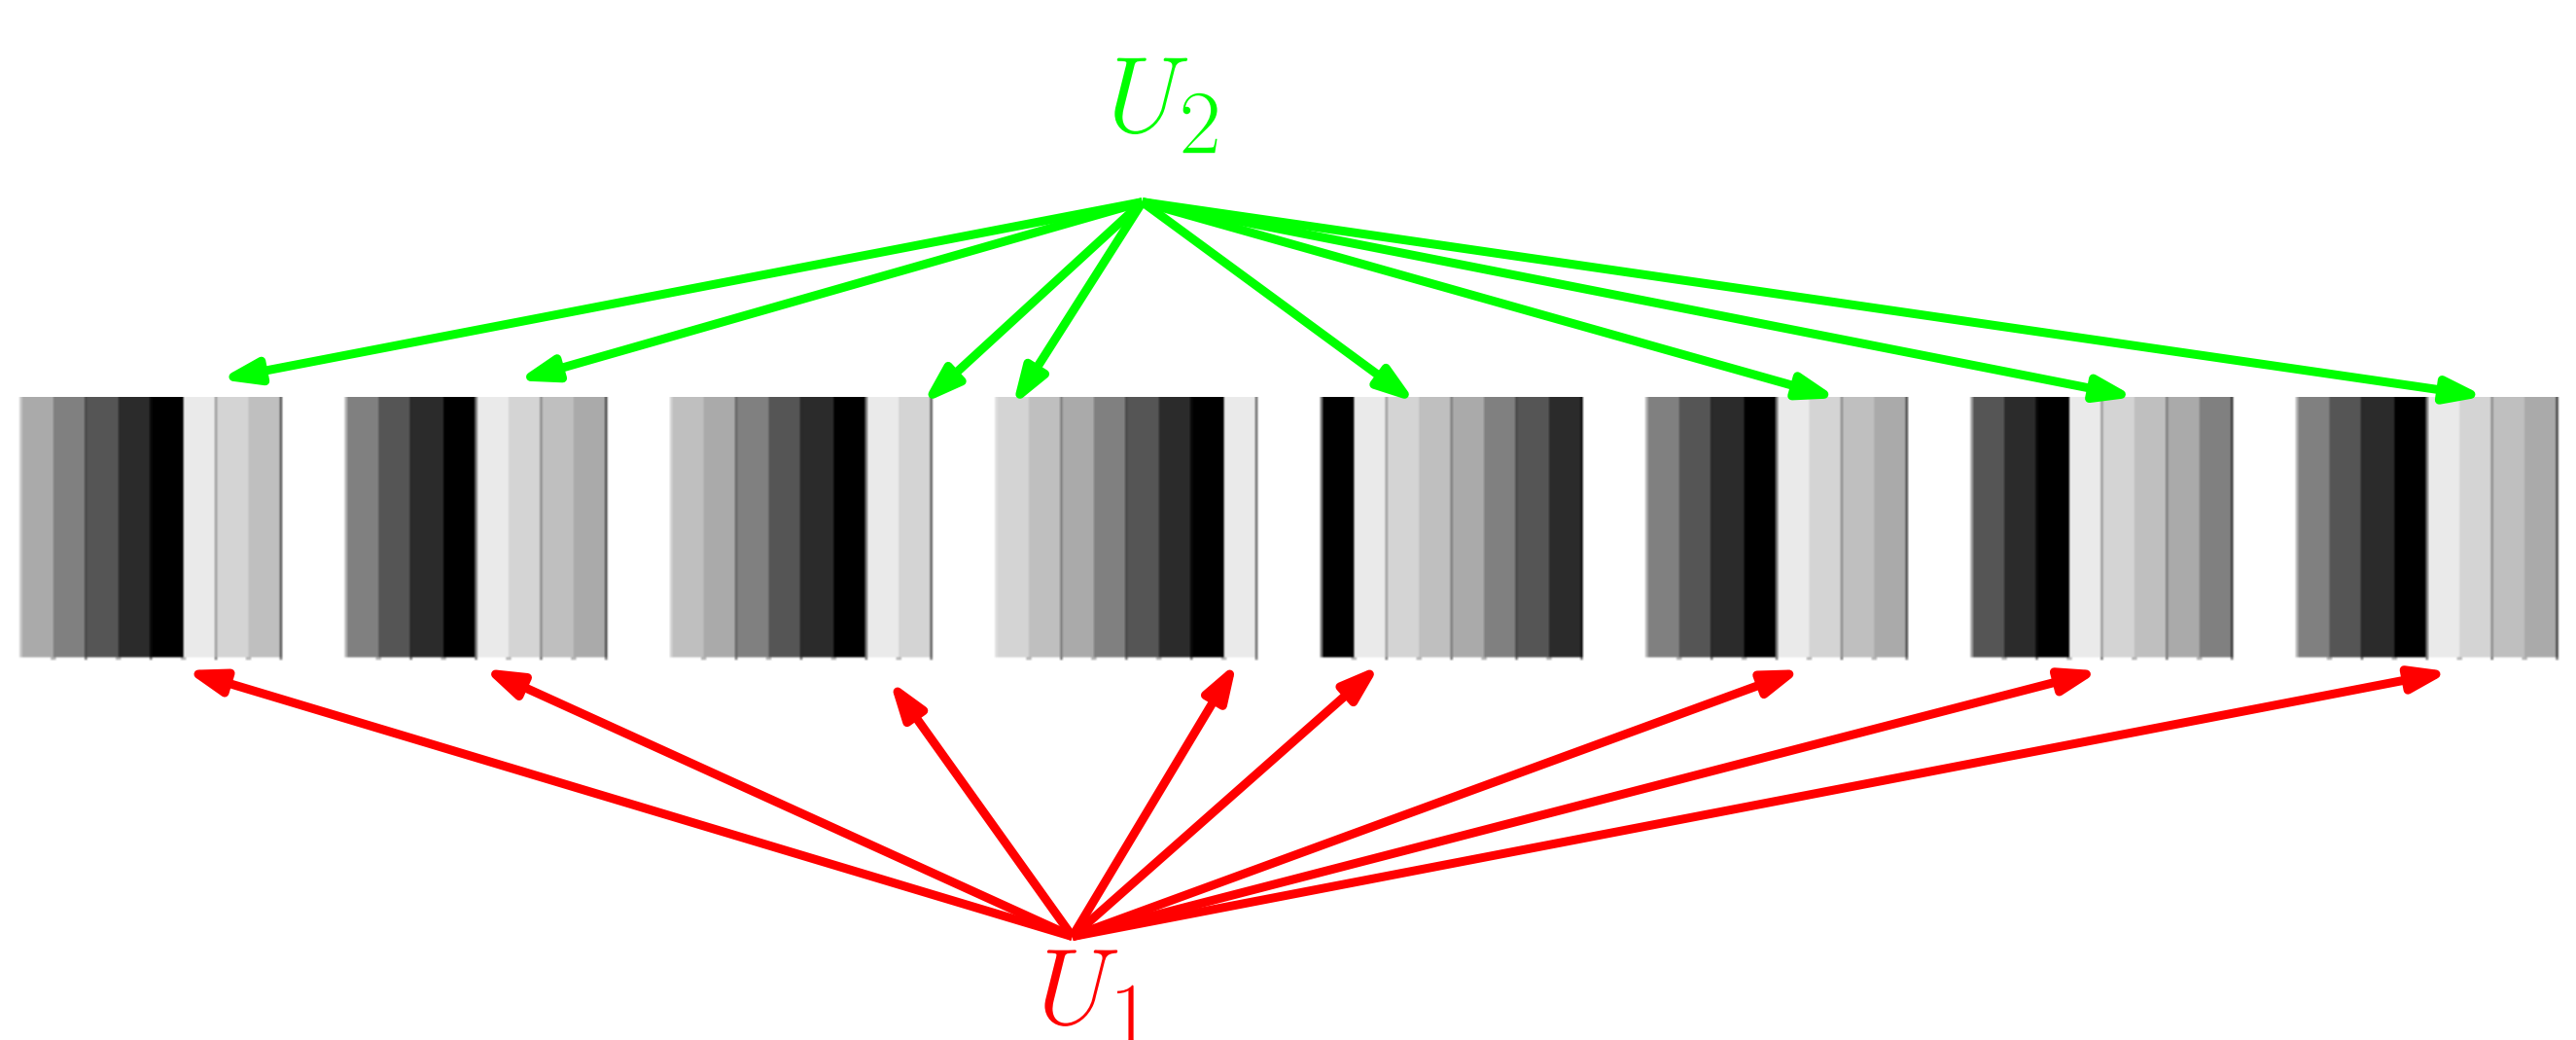
Smoothed-Striding Algorithm U_i .



HOW TO STORE THE GROUPS

Key Insight: While each U_i does need to contain a random element from each chunk, the U_i 's don’t need to be *independent*.

We store U_1 , and all other groups are determined by a “circular shift” of U_1 (wraparound within each chunk).



PARTIAL PARTITION STEP

The *Partial Partition Step* of the Smoothed Striding algorithm guarantees that all elements at index $i < v_{\min}$ have value $A[i] \leq \text{pivot value}$, and all elements at index $i > v_{\max}$ have value $A[i] > \text{pivot value}$. Thus, to completely partition the array a subarray of size $v_{\max} - v_{\min}$ must be partitioned. We prove the following proposition that bounds the size of this subarray:

Proposition:

Let $\epsilon \in (0, 1/2)$ and $\delta \in (0, 1/2)$ such that $\epsilon \geq \frac{1}{\text{poly}(n)}$ and $\delta \geq \frac{1}{\text{polylog}(n)}$. Suppose $s > \frac{\ln(n/\epsilon)}{\delta^2}$. Finally, suppose that each processor has a cache of size at least $s + c$ for a sufficiently large constant c .

Then the Partial-Partition Algorithm achieves work $O(n)$; achieves span $O(b \cdot s)$; incurs $\frac{s+c}{b} + O(1)$ cache misses; and guarantees with probability $1 - \epsilon$ that

$$v_{\max} - v_{\min} < 4n\delta.$$

RECURSIVE STRATEGIES

We propose two algorithms for solving the recursive subproblem:

- ▶ In the *Hybrid Smoothed Striding Algorithm* we recurse with a (Cache-Inefficient) In-Place Parallel-Partition algorithm, that has span $O(\log n \log \log n)$. With this recursive strategy we achieve span $O(\log n \log \log n)$ overall – which is within a $\log \log n$ factor of optimal – and incur fewer than $(n + o(n))/b$ cache misses – which is optimal up to low order terms– for appropriate parameter choices, with high probability in n .
- ▶ In the *Recursive Smoothed Striding Algorithm* we recurse with the Smoothed Striding algorithm. **This algorithm achieves span $O(\log^2 n)$ which is worse than the other approach**, but this algorithm **has the major benefit of simplicity to implement, while maintaining optimal cache behavior of $(n + o(n))/b$ for appropriate parameter choices, with high probability in n .**

ANALYSIS OVERVIEW

The proof of our proposition about the Parallel Partition Step proceeds along these lines:

- ▶ Let μ be the fraction of elements of the array that are less than the pivot, and μ_i be the fraction of elements of U_i that are less than the pivot.
- ▶ All the μ_i have identical probability distributions, because any given element of the array is equally likely to be assigned to any U_i . Hence $\mathbb{E}[\mu_i] = \mu$.
- ▶ $|U_i| = \text{polylog } n$, so a Hoeffding Bound (Chernoff Bound for random variable on $[0, 1]$ instead of on $\{0, 1\}$) guarantees that all U_i 's will have μ_i 's concentrated around μ with high probability in n .
- ▶ The concentration of μ_i 's induces a concentration of v_i 's.
- ▶ This guarantees that $v_{\max} - v_{\min}$ is small.

PSEUDOCODE FOR THE ALGORITHM

```

Recall:
A is the array to be partitioned, of length n.
We break A into chunks, each consisting of g cache lines of size b.
We create g groups U_1, ..., U_g that each contain a single cache line from each chunk.
U_i's j-th cache line is the (X[j] + i mod g + 1)-th cache line in the j-th chunk of A.

procedure GET BLOCK START INDEX(X, g, b, i, j)
    return b * ((X[j] + i mod g + 1) - 1) / g + 1
end procedure

procedure PARALLEL_PARTITION(A, n, g, b)
    if g < 2 then
        serial partition A
    else
        for j in {1, 2, ..., n/(gb)} do
            X[j] ← a random integer from [1, g]
        end for
        for all i in {1, 2, ..., g} in parallel do
            low ← GetBlockStartIndex(X, g, b, i, 1)
            high ← GetBlockStartIndex(X, g, b, i, n/(gb)) + b - 1
            while low < high do
                while A[low] ≤ pivotValue do
                    low ← low + 1
                    if low mod b = 0 then
                        low ← GetBlockStartIndex(X, g, b, i, k)
                    end if
                end while
                while A[high] > pivotValue do
                    high ← high - 1
                    if high mod b = 1 then
                        high ← GetBlockStartIndex(X, g, b, i, k') + b - 1
                    end if
                end while
                if low < high then
                    Swap A[low] and A[high]
                end if
            end while
            Recurse on A[low], ..., A[high - 1]
        end for
    end procedure
```