

RECURSIVE STRATEGIES

The smoothed striding algorithm allows for the subproblems to be solved with recursion. There are two algorithms that can be used in recursion:

- ▶ In the *Recursive Smoothed Striding Algorithm* we recurse with the smoothed striding algorithm. **This gets span $O(\log^2 n)$ which is not great**, but this algorithm **is very simple to implement, while maintaining optimal cache behavior**.
- ▶ We can also recurse with a Cache-Inneficient In-Place Parallel-Partition algorithm, that was developed concurrently to our algorithm. Doing so we achieve span $O(\log n \log \log n)$ and optimal cache behavior.

ANALYSIS OVERVIEW

Let μ be the faction of elements of the array that are less than the pivot, and μ_i be the fraction of elements of U_i that are less than the pivot.

- ▶ Each U_i has a random element from each chunk of the array, so each element of each U_i is randomly either greater than or less than the pivot, with probabilities $1 - \mu, \mu$.
- ▶ $|U_i| = \text{polylog } n$, so a Chernoff Bound guarantees that all U_i ’s will have μ_i ’s similar to $\mathbb{E}[\mu_i] = \mu$ with high probability in n .
- ▶ The concentration of μ_i ’s induces a concentration of v_i ’s.
- ▶ This guarantees that $v_{\max} - v_{\min}$ is small.

PSEUDOCODE

Figure: The Smoothed Striding Algorithm

Recall:
 A is the array to be partitioned, of length n .
We break A into chunks, each consisting of g cache lines of size b .
We create g groups U_1, \dots, U_g that each contain a single cache line from each chunk.
 U_i ’s j -th cache line is the $(X[j] + i \bmod g + 1)$ -th cache line in the j -th chunk of A .

procedure GET BLOCK START INDEX(X, g, b, i, j)
 return $b \cdot ((X[j] + i \bmod g) + (j - 1) \cdot g) + 1$
end procedure

procedure PARALLELPARTITION(A, n, g, b)
 if $g < 2$ **then**
 serial partition A
 else
 for $j \in \{1, 2, \dots, n/(gb)\}$ **do**
 $X[j] \leftarrow$ a random integer from $[1, g]$
 end for
 for all $i \in \{1, 2, \dots, g\}$ **in parallel** **do**
 $\text{low} \leftarrow \text{GetBlockStartIndex}(X, g, b, i, 1)$
 $\text{high} \leftarrow \text{GetBlockStartIndex}(X, g, b, i, n/(gb)) + b - 1$
 while $\text{low} < \text{high}$ **do**
 while $A[\text{low}] \leq \text{pivotValue}$ **do**
 $\text{low} \leftarrow \text{low} + 1$
 if $\text{low} \bmod b \equiv 0$ **then**
 $k \leftarrow$ number of block increments so far (including this one)
 $\text{low} \leftarrow \text{GetBlockStartIndex}(X, g, b, i, k)$
 end if
 end while
 while $A[\text{high}] > \text{pivotValue}$ **do**
 $\text{high} \leftarrow \text{high} - 1$
 if $\text{high} \bmod b \equiv 1$ **then**
 $k' \leftarrow$ number of block decrements so far (including this one)
 $k' \leftarrow n/(gb) - k'$
 $\text{high} \leftarrow \text{GetBlockStartIndex}(X, g, b, i, k') + b - 1$
 end if
 end while
 Swap $A[\text{low}]$ and $A[\text{high}]$
 end while
 end for
 Recurse on $A[v_{\min}], \dots, A[v_{\max} - 1]$
 end if
end procedure

▷ This procedure returns the index in A of the start of U_i ’s j -th block.

▷ We perform a serial partition on all U_i ’s in parallel

▷ $\text{low} \leftarrow$ index of the first element in U_i

▷ $\text{high} \leftarrow$ index of the last element in U_i

▷ Perform a block increment once low reaches the end of a block

▷ Increase low to start of block k of G_i

▷ Perform a block decrement once high reaches the beginning of a block

▷ Decrease high to end of block k' of G_i