

Cache Efficient Parallel Partition Algorithms

An In-Place Exclusive Read/Write Memory Algorithm

WHY IS THE PARTITION PROBLEM IMPORTANT?

The Partition Problem is a fundamental problem in computer science. Additionally, it is used in many algorithms such as

- Parallel Quicksort. This is the most well-known application of parallel-partition. Sorting is a very fundamental problem.
- Filtering operations.

Humans like sorted data.
Computers like sorted data.

OUR RESEARCH QUESTION


Can we create an algorithm with *theoretical guarantees* that is *fast in practice*?

RESULT

We created the *Smoothed Striding Algorithm*.
Key Features:

- linear work and polylogarithmic span (like the Standard Algorithm)
- fast in practice (like the Strided Algorithm)
- theoretically optimal cache behavior (unlike any past algorithm)

SMOOTHED STRIDING ALGORITHM’S PERFORMANCE



STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

Strided Algorithm

[Francis and Pannan, 92;
Frias and Petit, 08]

- Good cache behavior in practice
- Worst case span is $T_\infty \approx n$
- On random inputs span is $T_\infty = \tilde{O}(n^{2/3})$

Smoothed-Striding Algorithm

- Provably optimal cache behavior
- Span is $T_\infty = O(\log n \log \log n)$ with high probability in n
- Uses randomization *inside* the algorithm

APPLICATION TO PARALLEL QUICKSORT


Parallel Quicksort is the most important application of Parallel Partition. Parallel Quicksort works as follows:

- Chose a pivot value randomly from the array
- *Partition* the array relative to the pivot value
- Recursively sort the subarrays (in parallel)

The depth of recursion is $O(\log n)$ with high probability in n , and each level of recursion requires span $O(\log n \log \log n)$ when using the smoothed striding algorithm, which results in span $O(\log^2 n \log \log n)$ and work $O(n \log n)$ for the entire parallel quicksort algorithm, which is within a factor of $\log \log n$ of optimal, while additionally guaranteeing that the algorithm is cache-friendly.


SMOOTHED STRIDING ALGORITHM

Logically partition the array into chunks of adjacent elements.




Form groups U_i that contain a random element from each chunk.


This randomization step was one of our key insights; it guarantees that the U_i 's have similar compositions regardless of the input.




Perform serial partitions on each U_i in parallel over the U_i 's.
This step is highly parallel.



Define $v_i = \text{index of first element greater than the pivot in } U_i$.




Identify leftmost and rightmost v_i . Note that $A[k] \leq \text{pivot}$ for all $k < v_{\min}$, and $A[k] > \text{pivot}$ for all $k \geq v_{\max}$.



Recursively partition the subarray.

This step was previously impossible; adding randomization enables this step, which enables our algorithm's low span.




A KEY CHALLENGE


How do we store the U_i 's if they are all random?

Storing which elements make up each U_i takes too much space!

Strided Algorithm P_i .



Smoothed-Striding Algorithm U_i .



HOW TO STORE THE GROUPS

Key Insight: While each U_i does need to contain a random element from each chunk, the U_i 's don't need to be *independent*.

We store U_1 , and all other groups are determined by a “circular shift” of U_1 (wraparound within each chunk).

