

# Cache Efficient Parallel Partition Algorithms

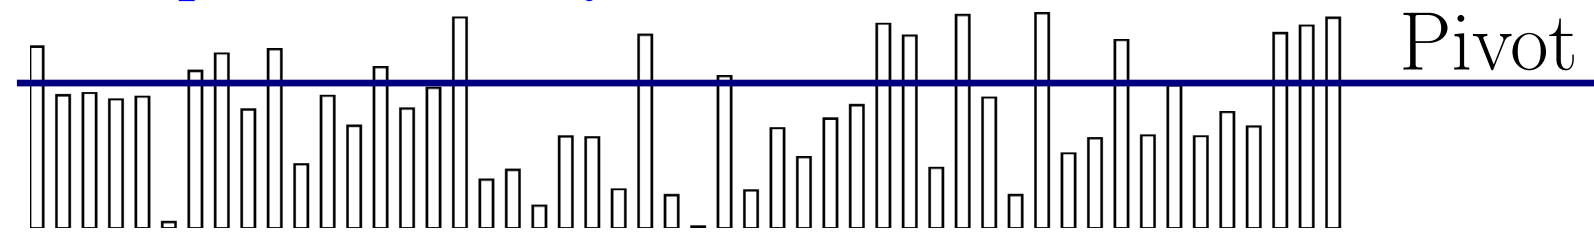
## An In-Place Exclusive Read/Write Memory Algorithm

### WHAT IS THE PARTITION PROBLEM?

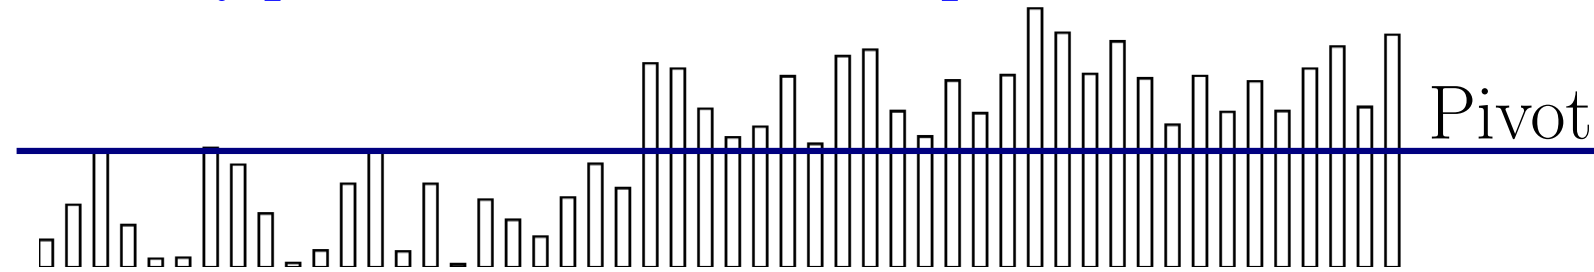
**Explanation:** The *Partition Problem* is to reorder the elements in a list so that elements in the same group occur in the same part of the list.

**Example:** A common way of grouping elements is based on whether they exceed or fall short of a certain “pivot” value.

An unpartitioned array:



An array partitioned relative to the pivot value:

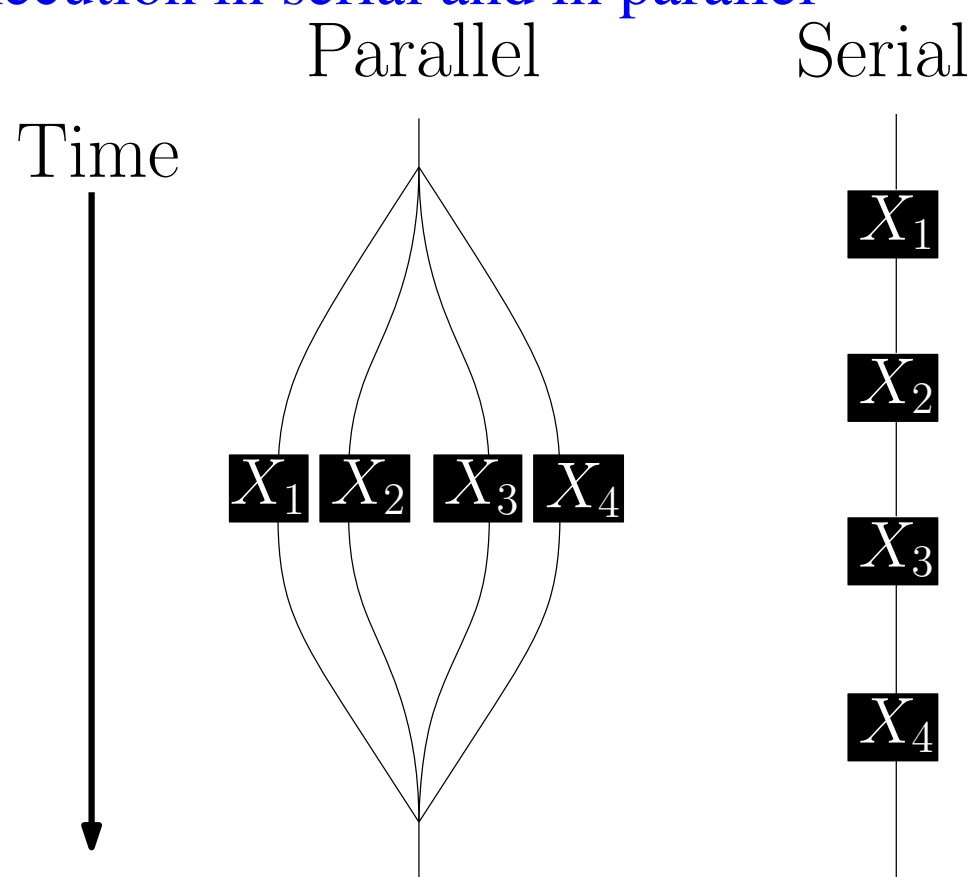


### WHAT IS A PARALLEL ALGORITHM?

**Explanation:** Whereas a typical (i.e. serial) algorithm runs on a single processor, a *parallel algorithm* runs on  $p \geq 1$  processors.

**Example:** Many tasks have parts that can be performed concurrently; such tasks can be performed faster with parallel computing.

Program execution in serial and in parallel



### WHAT IS CACHE EFFICIENCY?

**Explanation:** *Cache* is a small part of memory that can be accessed much faster than ordinary RAM. When data is already loaded into Cache a program can rapidly access it; this is called a *cache hit*. When data needed by a program isn't in cache it must be loaded into cache; this is called a *cache miss*, and takes time.

**Remark:** An algorithm with very few cache misses is *Cache Efficient*; cache efficiency leads to faster performance in practice.

**Factors in Cache-Efficiency:**

- Perform low number of passes over the data
- Don't use extra memory, i.e. are *In-Place*
- Deal with elements that are close in memory together

### PREVIOUS WORK ON THE PARTITION PROBLEM

The “Standard Algorithm” is *theoretically optimal with span  $O(\log n)$* , but *slow in practice due to poor cache behavior*. The *fastest algorithms in practice lack theoretical guarantees*

- Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

**Not Exclusive Read/Write Memory**

- The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

**No locks or atomic-variables, but no bound on span**

### OUR RESEARCH QUESTION

Can we create an algorithm with *theoretical guarantees* that is *fast in practice*?

### RESULT

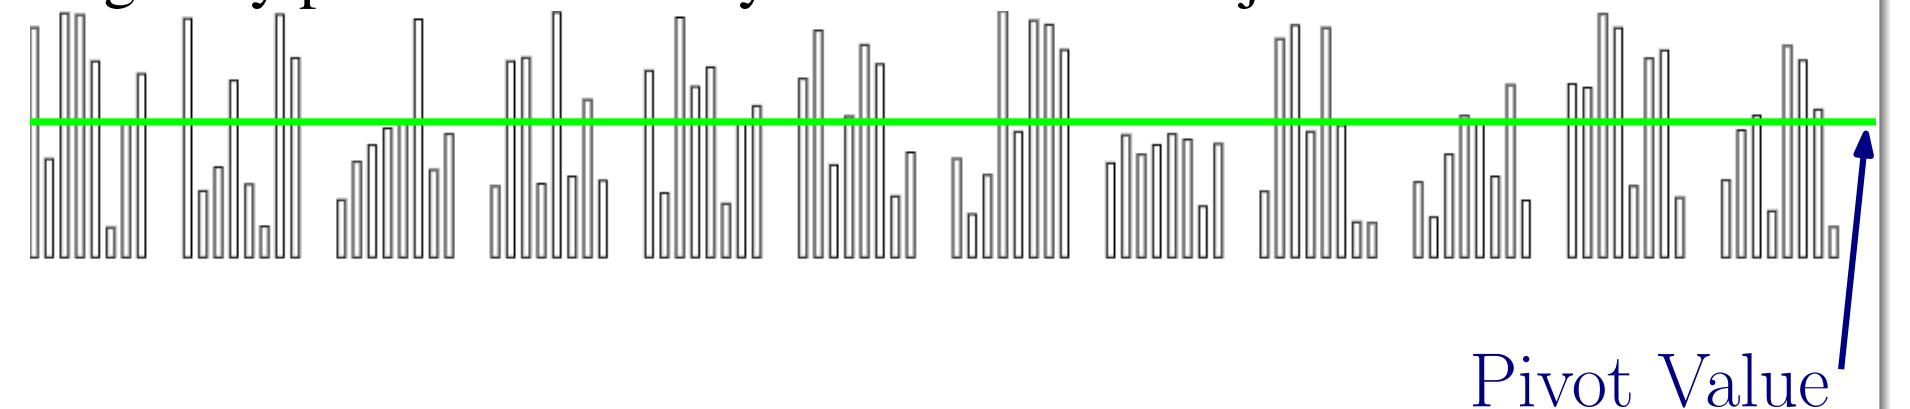
We created the *Smoothed Striding Algorithm*.

Key Features:

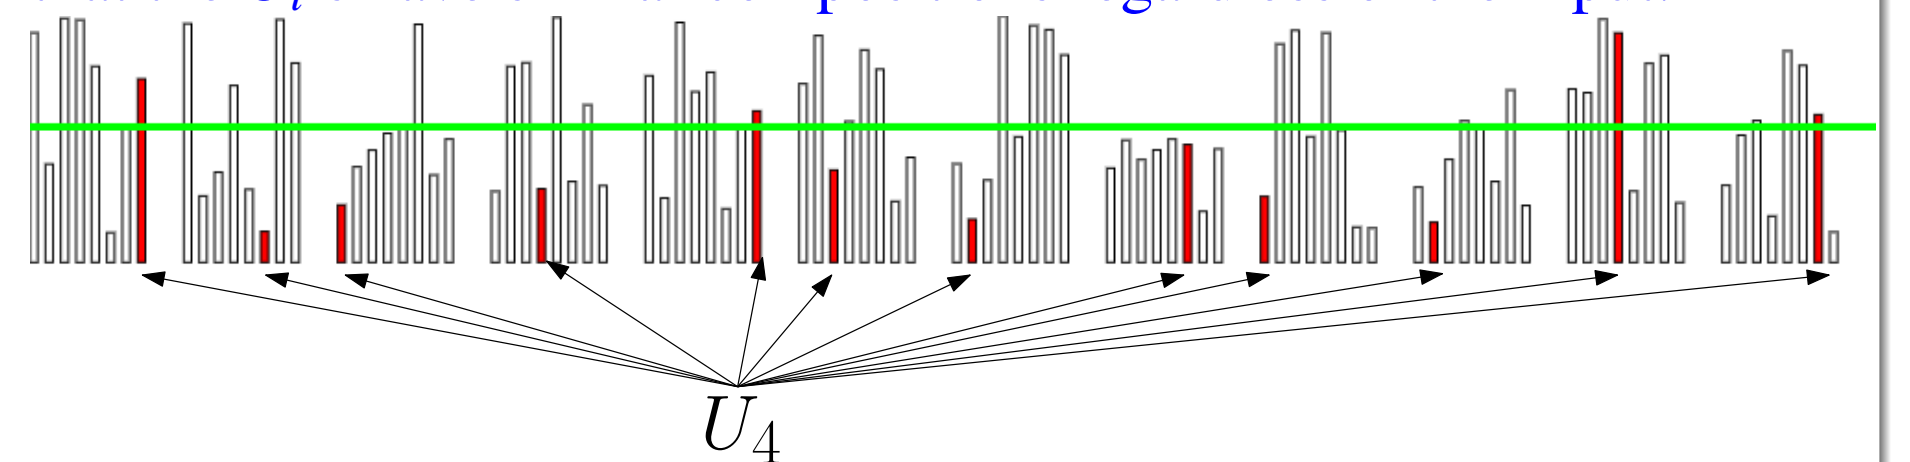
- linear work and polylogarithmic span  
(like the Standard Algorithm)
- fast in practice  
(like the Strided Algorithm)
- theoretically optimal cache behavior  
(unlike any past algorithm)

### SMOOTHED STRIDING ALGORITHM

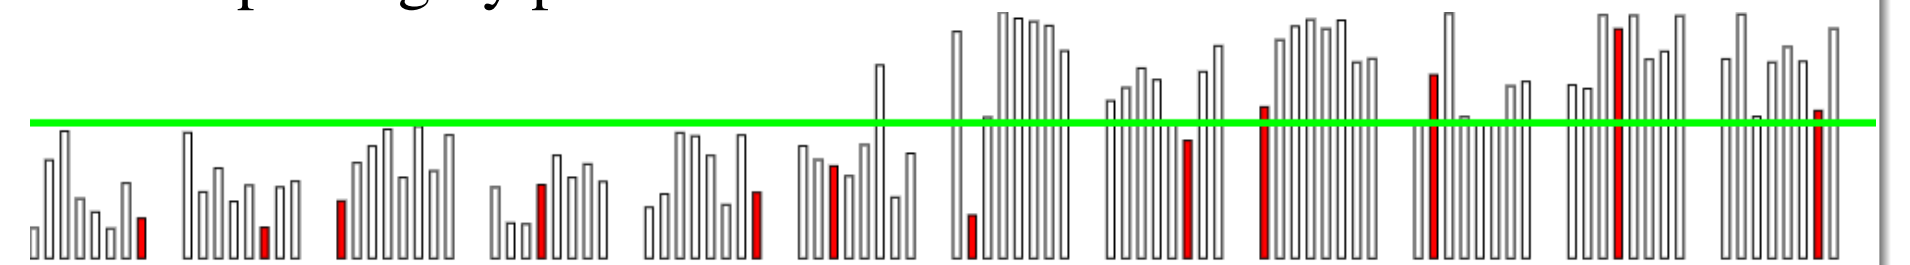
Logically partition the array into chunks of adjacent elements.



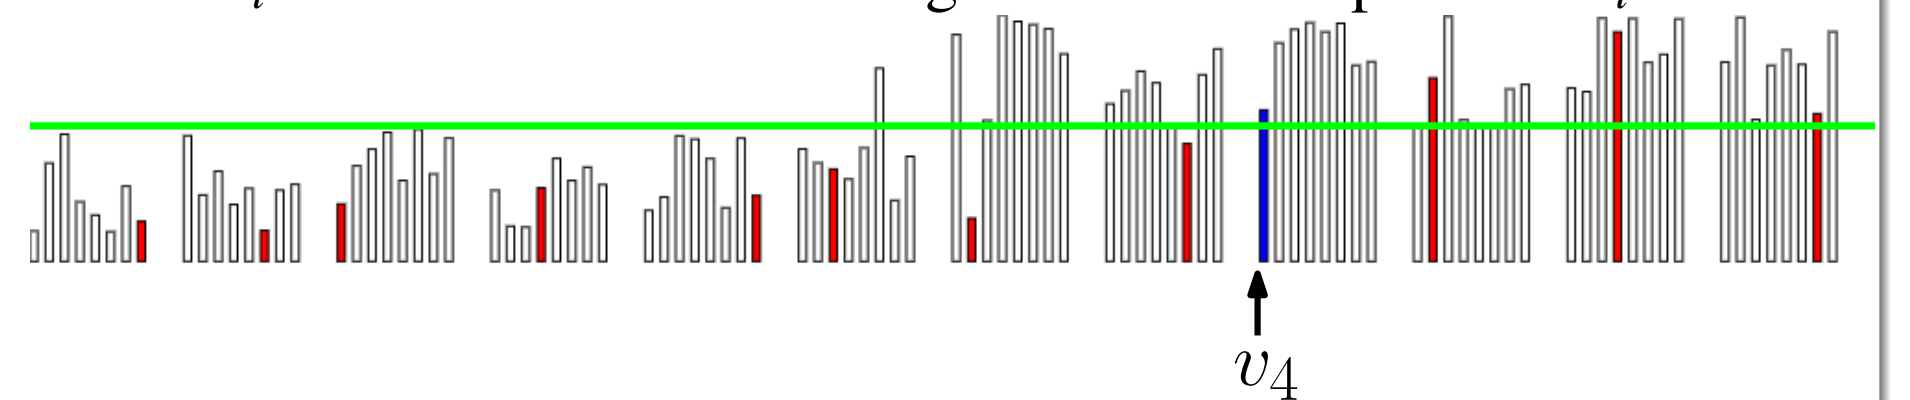
Form groups  $U_i$  that contain a random element from each chunk. This randomization step was one of our key insights; it guarantees that the  $U_i$ 's have similar compositions regardless of the input.



Perform serial partitions on each  $U_i$  in parallel over the  $U_i$ 's. This step is highly parallel.

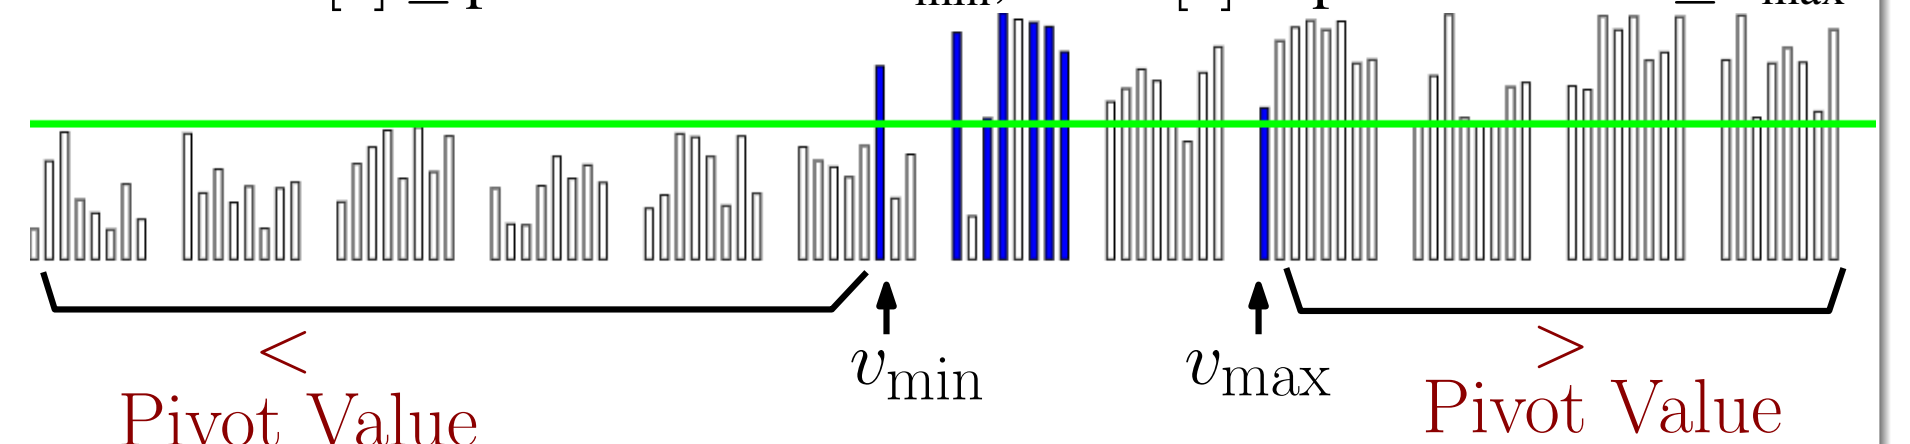


Define  $v_i$  = index of first element greater than the pivot in  $U_i$ .



Identify leftmost and rightmost  $v_i$ .

Note that  $A[k] \leq \text{pivot}$  for all  $k < v_{\min}$ , and  $A[k] > \text{pivot}$  for all  $k \geq v_{\max}$ .



Recursively partition the subarray.

This step was previously impossible; adding randomization enables this step, which enables our algorithm's low span.

