

Cache-Efficient Parallel Partition Algorithms

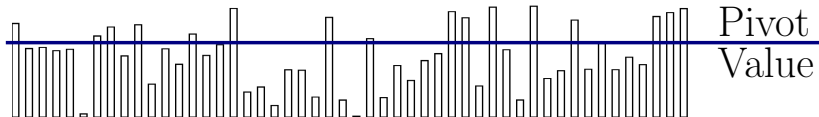
Alek Westover

MIT PRIMES

October 20, 2019

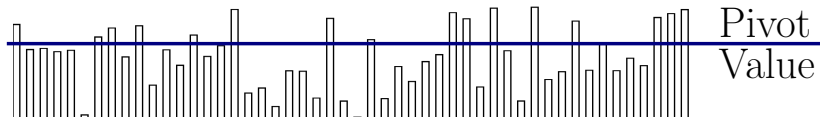
THE PARTITION PROBLEM

An unpartitioned array:

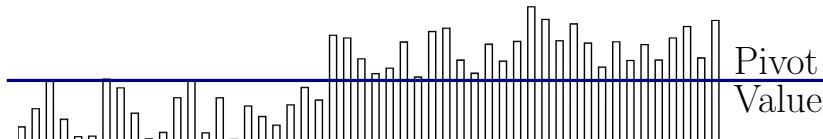


THE PARTITION PROBLEM

An unpartitioned array:



An array partitioned relative to a pivot value:



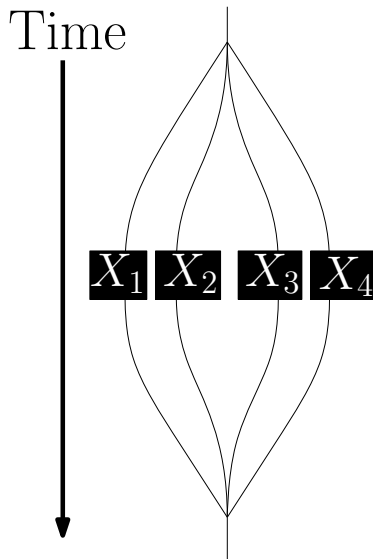
WHAT IS A PARALLEL ALGORITHM?

Fundamental primitive:

Parallel for loop


Parallel-For i from 1 to 4:

Do X_i

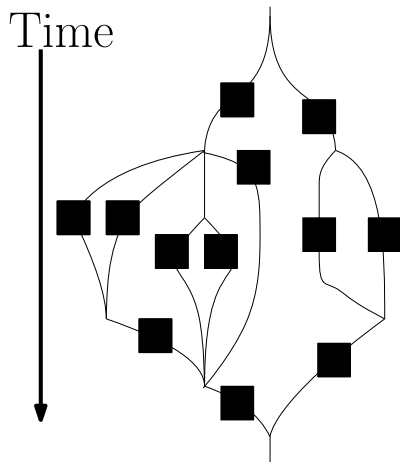


WHAT IS A PARALLEL ALGORITHM?

More complicated parallel structures can be made by combining parallel for loops and recursion.



T_p : TIME TO RUN ON p PROCESSORS



Important extreme cases:

Work: T_1 ,

- ▶ time to run in serial
- ▶ "sum of all work"

Span: T_∞ ,

- ▶ time to run on infinitely many processors,
- ▶ "height of the graph"

BOUNDING T_p WITH WORK AND SPAN

Brent's Theorem: [Brent, 74]

$$T_p = \Theta \left(\frac{T_1}{p} + T_\infty \right)$$

Take away: Work T_1 and span T_∞ determine T_p .

THE STANDARD PARALLEL PARTITION ALGORITHM

<i>Step</i>	<i>Span</i>
Create filtered array	$O(1)$
Compute prefix sums of filtered array	$O(\log n)$
Use prefix sums to partition array	$O(1)$

Total span: $O(\log n)$

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
- ▶ Makes multiple passes over array

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
 - ▶ Makes multiple passes over array
- } "bad cache behavior"

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
 - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
 - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippos Tsigas and Yi Zhang, 2003]

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
 - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
 - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

No locks or atomic-variables, but no bound on span

THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
 - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

No locks or atomic-variables, but no bound on span

Our Question: Can we create an algorithm with theoretical guarantees that is fast in practice?

OUR RESULT: THE SMOOTHED-STRIDING ALGORITHM

The Smoothed-Striding algorithm:

- ▶ has linear work and polylogarithmic span
(like the Standard Algorithm)
- ▶ is fast in practice
(like the Strided Algorithm)
- ▶ has theoretically optimal cache behavior
(unlike any past algorithm)

STRIDED VERSUS SMOOTHED STRIDED ALGORITHM

Strided Algorithm


[Francis and Pannan, 92; Frias and Petit, 08]

- ▶ Good cache behavior in practice
- ▶ Worst case span is $T_\infty \approx n$
- ▶ On random inputs span is $T_\infty = \tilde{O}(n^{2/3})$

Smoothed-Striding Algorithm

- ▶ Provably optimal cache behavior
- ▶ Span is $T_\infty = O(\log n \log \log n)$ with high probability in n
- ▶ Uses randomization *inside* the algorithm


SPEEDUP OVER SERIAL PARTITION: SMOOTHED-STRIDING ALGORITHM'S PERFORMANCE




The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]


Logically partition the array into chunks of adjacent elements:



Form groups P_i that contain the i -th element from each chunk:




Perform serial partitions on each P_i in parallel over the P_i 's:




This step is highly parallel.


Define $v_i = \text{index of first element greater than the pivot in } P_i$.




Identify leftmost and rightmost v_i .



Final step: Recursively partition the subarray




Final step: Recursively partition the subarray



- ▶ Recursion is impossible!
- ▶ **Final Step:** Partition the subarray *in serial*.

Subproblem Span $T_{\infty} \approx v_{\max} - v_{\min}$

~~Final step: Recursively partition the subarray~~




- ▶ Recursion is impossible!
- ▶ **Final Step:** Partition the subarray *in serial*.


Subproblem Span $T_{\infty} \approx v_{\max} - v_{\min}$ \leftarrow n in worst case.

The Smoothed-Striding Algorithm


Logically partition the array into chunks of adjacent elements:



Key difference: Form groups U_i that contain a random element from each chunk (rather than containing the i -th element from each chunk every time as the Strided Algorithm's P_i 's do)




Perform serial partitions on each U_i in parallel over the U_i 's:




This step is highly parallel.


Define $v_i = \text{index of first element greater than the pivot in } U_i$.




Identify leftmost and rightmost v_i .



Final step: Recursively partition the subarray



Final step: Recursively partition the subarray



► Recursion is now possible!

Subproblem can be solved with an In-Place algorithm that has span $O(\log n \log \log n)$.

A KEY CHALLENGE


How do we store the U_i 's if they are all random?

Storing which elements make up each U_i takes too much space!

Strided Algorithm P_i .



Smoothed-Striding Algorithm U_i .



AN OPEN QUESTION

Our algorithm: $\text{span } T_\infty = O(\log n \log \log n)$

Standard Algorithm: $\text{span } T_\infty = O(\log n)$.

Can we get optimal cache behavior and $\text{span } O(\log n)$?

ACKNOWLEDGMENTS

I would like to thank

- ▶ MIT PRIMES
- ▶ William Kuszmaul, my PRIMES mentor
- ▶ My parents