

# Cache-Efficient Parallel Partition Algorithms


Alek Westover

MIT PRIMES

October 20, 2019

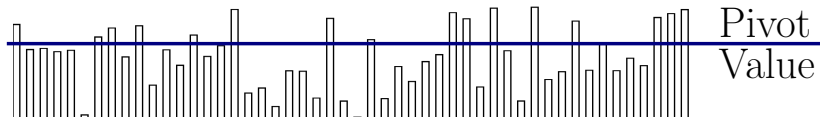
# THE PARTITION PROBLEM

An unpartitioned array:

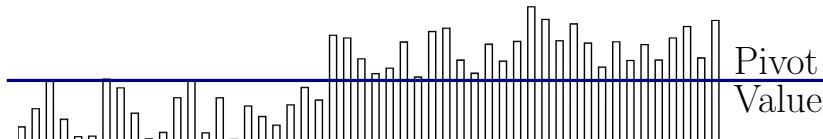


# THE PARTITION PROBLEM

An unpartitioned array:



An array partitioned relative to a pivot value:



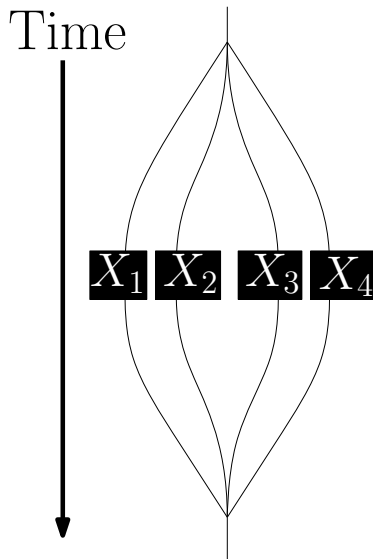
# WHAT IS A PARALLEL ALGORITHM?

Fundamental primitive:

*Parallel for loop*


**Parallel-For**  $i$  from 1 to 4:

Do  $X_i$

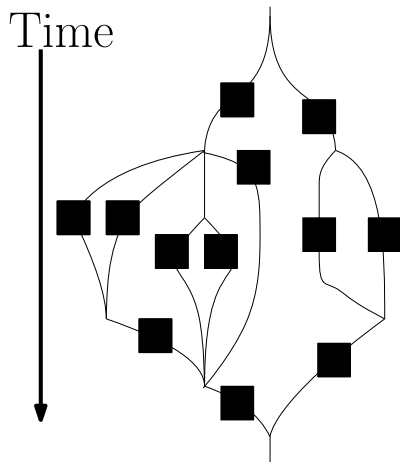


# WHAT IS A PARALLEL ALGORITHM?

More complicated parallel structures can be made by combining parallel for loops and recursion.



## $T_p$ : TIME TO RUN ON $p$ PROCESSORS



Important extreme cases:

**Work:**  $T_1$ ,

- ▶ time to run in serial
- ▶ "sum of all work"

**Span:**  $T_\infty$ ,

- ▶ time to run on infinitely many processors,
- ▶ "height of the graph"

# BOUNDING $T_p$ WITH WORK AND SPAN

**Brent's Theorem:** [Brent, 74]

$$T_p = \Theta \left( \frac{T_1}{p} + T_\infty \right)$$

**Take away:** Work  $T_1$  and span  $T_\infty$  determine  $T_p$ .

# THE STANDARD PARALLEL PARTITION ALGORITHM

<i>Step</i>	<i>Span</i>
Create filtered array	$O(1)$
Compute prefix sums of filtered array	$O(\log n)$
Use prefix sums to partition array	$O(1)$

Total span:  $O(\log n)$



# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
- ▶ Makes multiple passes over array

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
  - ▶ Makes multiple passes over array
- } "bad cache behavior"

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
  - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
  - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippos Tsigas and Yi Zhang, 2003]

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
  - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
  - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

No locks or atomic-variables,  
but no bound on span in general



# THE PROBLEM

Why is the Standard Algorithm is slow in practice?

- ▶ Uses extra memory
  - ▶ Makes multiple passes over array
- } "bad cache behavior"

But fastest algorithms in practice lack theoretical guarantees

- ▶ Lock-based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

No locks or atomic-variables,  
but no bound on span in general

**Our Question:** Can we create an algorithm with theoretical guarantees that is fast in practice?

# OUR RESULT

We created a randomized algorithm for the parallel partition problem: the *Smoothed-Striding Algorithm*.

The Smoothed-Striding algorithm has...

- ▶ performance comparable to that of the Strided Algorithm
- ▶ polylogarithmic span like the Standard Algorithm
- ▶ theoretically optimal cache behavior

# ALGORITHM COMPARISON

## Strided Algorithm

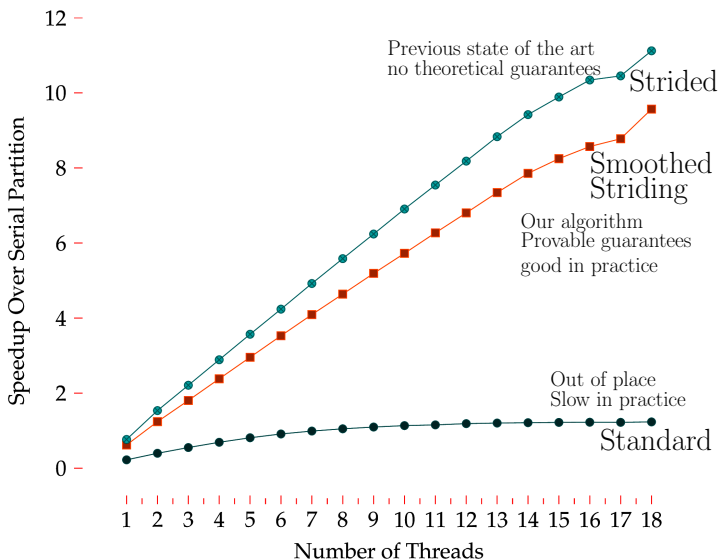
[Francis and Pannan, 92; Frias and Petit, 08]

- ▶ Good cache behavior in practice
- ▶ Worst case span is  $T_\infty \approx n$
- ▶ But: on random inputs span is  $T_\infty = \tilde{O}(n^{2/3})$

## Smoothed-Striding Algorithm

- ▶ Provably optimal cache behavior
- ▶ Span is  $O(\log n \log \log n)$  with high probability in  $n$
- ▶ Uses randomization inside the algorithm to remove the need for randomized input

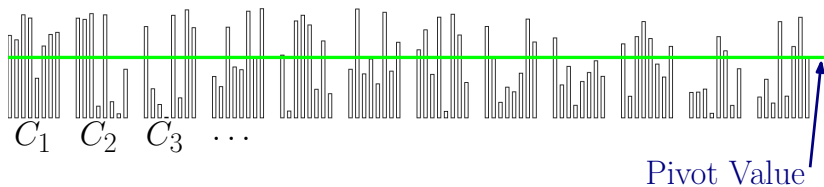
# SPEEDUP OVER SERIAL PARTITION: SMOOTHED-STRIDING ALGORITHM'S PERFORMANCE



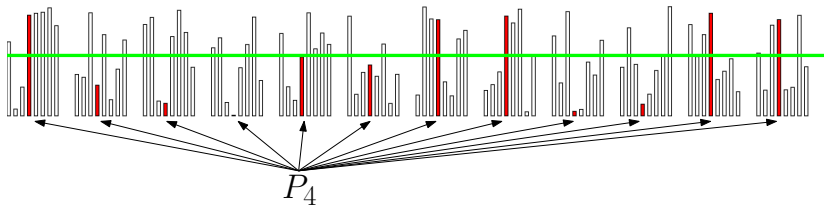
# The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

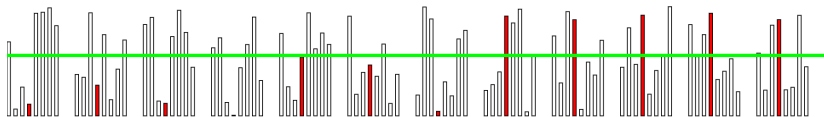
Logically partition the array into chunks of adjacent elements:



Form groups  $P_i$  where  $P_i$  contains the  $i$ -th element from each chunk:

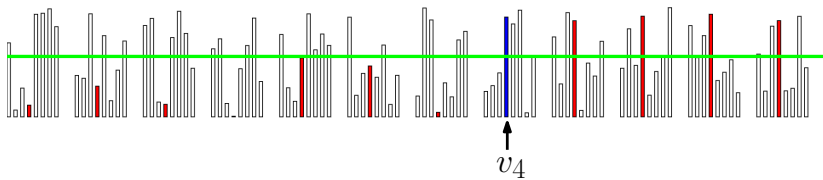


Perform serial partitions on each  $P_i$  in parallel over the  $P_i$ 's:

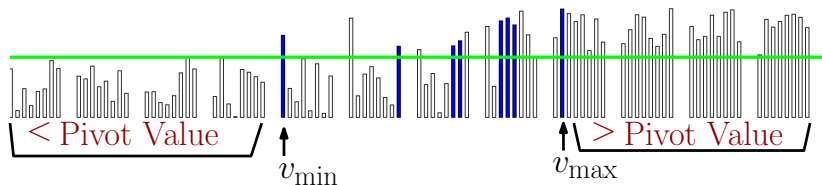


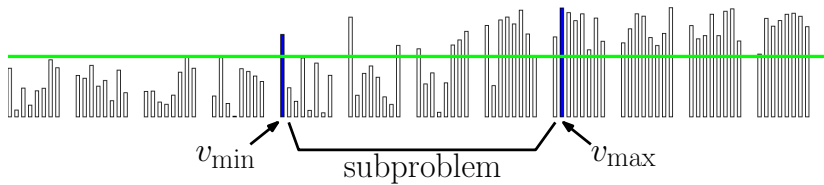
This step is highly parallel.



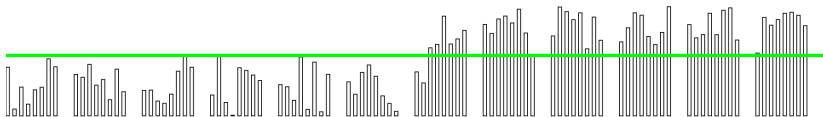


Identify the splitting index  $v_i$  (the first element greater than the pivot) of each  $P_i$ .





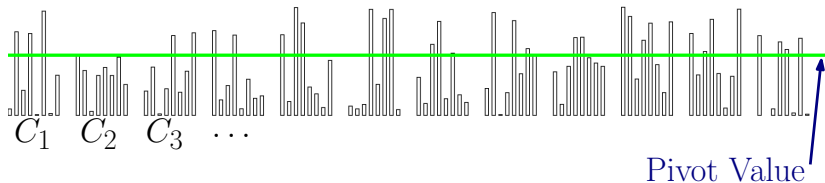
Partition the subarray from the minimum splitting index to the maximum splitting index in serial. This completes the partition.



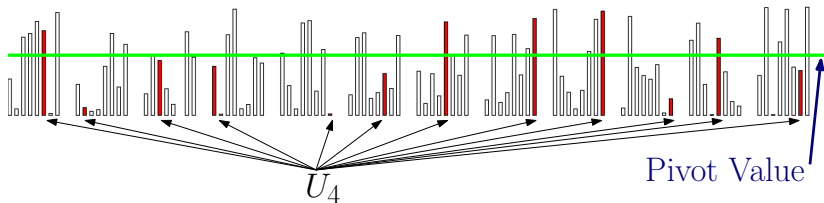
Note that this step has no parallelism. In general this results in span  $O(n)$ . However, if the number of elements less than the pivot in each  $P_i$  is similar, then size of the subarray to be partitioned can be very small.

# The Smoothed-Striding Algorithm

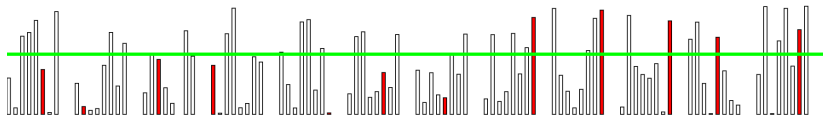
Logically partition the array into chunks of adjacent elements:



Form groups  $U_i$  where  $U_i$  contains the  $i$ -th element from each chunk:



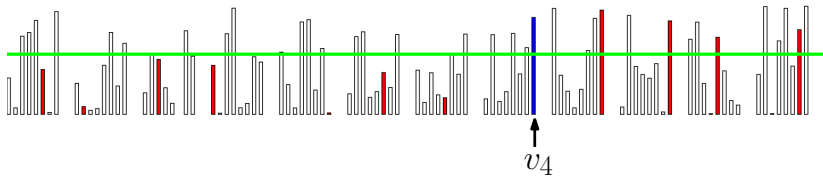
Perform serial partitions on each  $U_i$  in parallel over the  $U_i$ 's:




This step is highly parallel.

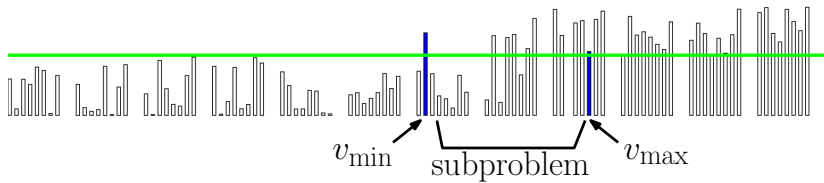


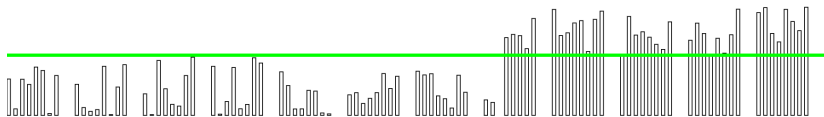
Identify the splitting index  $v_i$  (the first element greater than the pivot) of each  $P_i$ .



Recursively apply this algorithm to partition the subarray from the minimum splitting index to the maximum splitting index in serial. This completes the partition.








Unlike in the Strided Algorithm this step has parallelism, and is guaranteed to only run on a small subarray. The Strided Algorithm could not recurse here because the subproblem is a worst case input for it.

## A KEY CHALLENGE


How do we store the  $U_i$ 's if they are all random?

If we stored the parts that made up each  $U_i$  we would not be in-place.

**Blocked Strided Algorithm  $P_i$ .**



**Smoothed-Striding Algorithm  $U_i$ .**



## OPEN QUESTIONS

By recursively applying the Smoothed-Striding algorithm we get an algorithm for parallel partition that incurs  $n(1 + o(1))$  cache misses and has span  $O(b \log^2 n)$ .

There are techniques for improving this span to  $O(\log n \log \log n)$  while retaining the cache behavior.

But the standard algorithm has span  $O(\log n)$ .

Can we construct an algorithm that achieves optimal cache behavior and span  $O(\log n)$ ?

# ACKNOWLEDGMENTS

I would like to thank

- ▶ The MIT PRIMES program
- ▶ William Kuszmaul, my PRIMES mentor
- ▶ My parents