

Cache-Efficient Parallel Partition Algorithms

Alek Westover

MIT PRIMES

October 20, 2019

Partition

Definition (Partitioned Array)

$A[i]$ predecessor, $A[j]$ successor $\implies i < j$

Definition (Array partitioned relative to pivot value p)

$A[i] \leq p, A[j] > p \implies i < j$

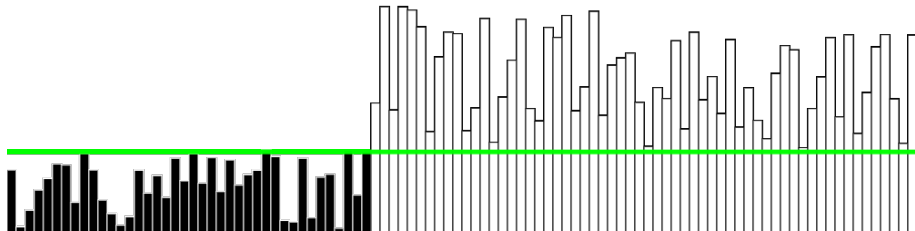
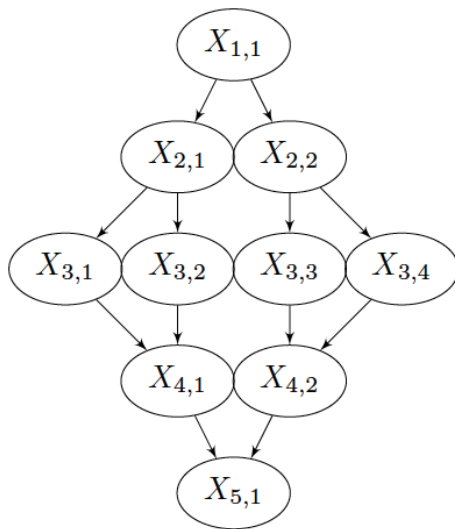


Figure: Black: Predecessor, White: Successor, Green: Pivot Value

Parallel Algorithm



Work and Span

Definition (T_p)

Running time on p processors. Note: $T_p \geq T_\infty$, $T_p \geq \frac{T_1}{p}$.

Definition (Work)

Running time on a single processor. $T_1 = \sum_i W_i$.

Definition (Span)

Running time on infinitely many processors. $T_\infty = \max_i W_i$.

Theorem (Brent's Theorem)

$$T_p = \sum_i \left\lceil \frac{W_i}{p} \right\rceil \leq \sum_i \left(\frac{W_i}{p} + 1 \right) = \frac{T_1}{p} + T_\infty.$$

Serial Partition

```
while low < high do  
    while  $A[\text{low}] \leq \text{pivotValue}$  do  
        low  $\leftarrow$  low + 1  
    end while  
    while  $A[\text{high}] > \text{pivotValue}$  do  
        high  $\leftarrow$  high - 1  
    end while  
    Swap  $A[\text{low}]$  with  $A[\text{high}]$   
end while  
if  $A[\text{low}] \leq \text{pivotValue}$  then  
    low  $\leftarrow$  low + 1  
end if
```

Randomized Algorithms

Definition (With high probability in n)

Probability of success is

$$1 - \frac{1}{n^c}$$

for c of our choice. i.e. the probability can be made arbitrarily close to 1.

Memory Bandwidth Bound

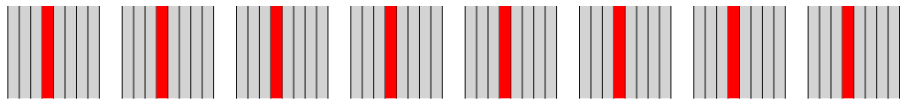
Definition (Cache miss)

A cache miss occurs when the algorithm must load a cache-line not stored in cache into cache.

- Being in-place is desirable
- Kuszmaul developed an in-place version of parallel partition.
- His in-place version outperformed the out-of-place standard algorithm.
- But was outperformed in practice by a lower span algorithm
- Experiments show the source of the problem: Memory bandwidth bound, i.e. incurring too many cache misses
- Note that both temporal and spatial cache efficiency are important.

Strided Algorithm citeFrancisPa92, Frias08 Description

Partition A into chunks $C_1, C_2, \dots, C_{n/gb}$ each consisting of g cache lines of size b . Let P_i be the union of the i -th cache-line from each chunk C_j .



Definition (Partially Partitioned Array)

$\exists u, l$ such that

$$i < u \implies A[i] \text{ is predecessor}, i \geq l \implies A[i] \text{ is successor}$$

- Perform serial partitions on all P_i in parallel.
- Let v_i be the position in A of the first successor in P_i . Perform a serial partition on $A[\min_i v_i], \dots, A[\max_i v_i - 1]$.

Strided Algorithm Analysis

- Partial partition step: work $O(n)$, span $\Theta(n/g)$.
- Serial cleanup step: span $\Theta(v_{\max} - v_{\min})$, which is $O(n)$ in general.
- If the number of predecessors in each P_i is similar, $v_{\max} - v_{\min}$ can be small.
- In particular, if $b \in \text{polylog}(n)$, and the array values are selected independently at random from some distribution, and g is chosen to optimize span ($g = n^{1/3}$), then with high probability in n ,

$$v_{\max} - v_{\min} < \tilde{O}(n^{2/3}),$$

the span is

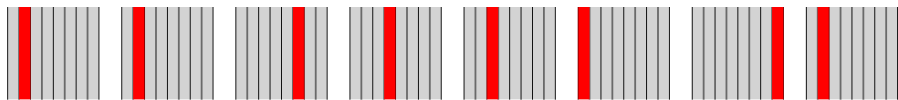
$$\tilde{O}(n^{2/3}),$$

and the number of cache misses is fewer than

$$\frac{n}{b} + \frac{\tilde{O}(n^{2/3})}{b}.$$

Smoothed Striding Algorithm Description

Let $X[1], \dots, X[s]$ be chosen uniformly at random from $\{1, \dots, g\}$. Let U_i be the union of the $(X[j] + i) \bmod g$ -th cache-line from each chunk C_j .



- Perform serial partitions on all U_i in parallel.
- The array is partially now partitioned with $A[i]$ a predecessor for all $i < v_{\min}$ and $A[i]$ a successor for all $i \geq v_{\max}$.

Note that we will make $s = \frac{n}{gb} < \text{polylog}(n)$ so the algorithm remains in-place.

Partial Partition Step Analysis

Proposition

Let $\epsilon \in (0, 1/2)$ and $\delta \in (0, 1/2)$ such that $\epsilon \geq \frac{1}{\text{poly}(n)}$ and $\delta \geq \frac{1}{\text{polylog}(n)}$.

Suppose $s > \frac{\ln(n/\epsilon)}{\delta^2}$. Finally, suppose that each processor has a cache of size at least $s + c$ for a sufficiently large constant c .

Then the Partial-Partition Algorithm achieves work $O(n)$; achieves span $O(b \cdot s)$; incurs $\frac{s+n}{b} + O(1)$ cache misses; and guarantees with probability $1 - \epsilon$ that

$$v_{\max} - v_{\min} < 4n\delta.$$

Hybrid Algorithm Analysis - General Theorem

Theorem

The Hybrid Smoothed Striding Algorithm algorithm using parameter $\delta \in (0, 1/2)$ satisfying $\delta \geq 1/\text{polylog}(n)$: has work $O(n)$; achieves span

$$O\left(\log n \log \log n + \frac{b \log n}{\delta^2}\right),$$

with high probability in n ; and incurs fewer than

$$(n + O(n\delta))/b$$

cache misses with high probability in n .

Hybrid Algorithm Analysis - Corollary for specific parameter settings

An interesting corollary of the above theorem concerns what happens when b is small (e.g., constant) and we choose δ to optimize span.

Corollary

Suppose $b \leq o(\log \log n)$. Then the Cache-Efficient Full-Partition Algorithm algorithm using $\delta = \Theta(\sqrt{b / \log \log n})$, achieves work $O(n)$, and with high probability in n , achieves span $O(\log n \log \log n)$ and incurs fewer than $(n + o(n))/b$ cache misses.

Recursive Algorithm Analysis - General Theorem

Theorem

With high probability in n , the Recursive Smoothed Striding algorithm using parameter $\delta \in (0, 1/2)$ satisfying $\delta \geq 1/\text{polylog}(n)$: achieves work $O(n)$, attains span

$$O\left(b\left(\log^2 n + \frac{\log n}{\delta^2}\right)\right),$$

and incurs $(n + O(n\delta))/b$ cache misses.

Recursive Algorithm Analysis - Corollary for specific parameter settings

A particularly natural parameter setting for the Recursive algorithm occurs at $\delta = 1/\sqrt{\log n}$.

Corollary

With high probability in n , the Recursive Smoothed Striding Algorithm using parameter $\delta = 1/\sqrt{\log n}$: achieves work $O(n)$, attains span $O(b \log^2 n)$, and incurs $n/b \cdot (1 + O(1/\sqrt{\log n}))$ cache misses.

Experiments

- Strided Algorithm vs Smoothed Striding algorithm
- Cache misses

Acknowledgments

I would like to thank

- The MIT PRIMES program
- William Kuszmaul, my PRIMES mentor
- My parents

References



Rhys S. Francis and LJH Pannan. 1992.

A parallel partition for enhanced parallel quicksort.

Parallel Comput. 18, 5 (1992), 543–550.