

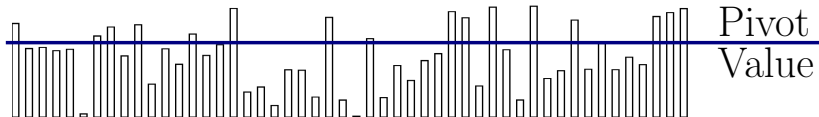
# Cache-Efficient Parallel Partition Algorithms Using Exclusive-Read-and-Write Memory

William Kuszmaul, Alek Westover

July 1, 2020

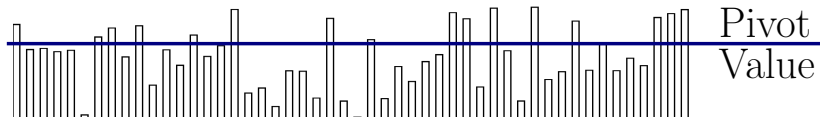
# THE PARTITION PROBLEM

An unpartitioned array:




# THE PARTITION PROBLEM

An unpartitioned array:



An array partitioned relative to a pivot value:




## EXAMPLE APPLICATIONS

- ▶ Parallel Partition
- ▶ Parallel Quicksort
- ▶ Filtering Operations


# WHAT IS A PARALLEL ALGORITHM?

Fundamental primitive: *Parallel for loop*



# PARALLEL ALGORITHM PERFORMANCE METRICS

Time



$T_p$ : Time to run on  $p$  processors

**Work:**  $T_1$

► time to run in serial

**Span:**  $T_\infty$

► time to run on infinitely many processors

Brent's Theorem:

$$T_p = \Theta \left( \frac{T_1}{p} + T_\infty \right)$$

# WHAT IS CACHE EFFICIENCY?

Roughly, an algorithm is **cache-efficient** if it

- ▶ Performs low number of passes over the data
- ▶ Doesn't use extra memory (i.e. is *In-Place*)
- ▶ Operates on elements close together in memory at the same time

Cache-Efficiency makes an algorithm faster in practice.

# THE PROBLEM

Standard Algorithm  $\text{span } O(\log n)$  but is **slow in practice**

- ▶ **Uses extra memory**
  - ▶ **Makes multiple passes over array**
- } bad cache behavior

**Fastest algorithms in practice lack theoretical guarantees**

- ▶ **Lock based** and **atomic-variable based** algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ **The Strided Algorithm**

[Francis and Pannan, 92; Frias and Petit, 08]

**No locks or atomic-variables, but no bound on span**



## OUR QUESTION

Can we create an algorithm with *theoretical guarantees* that is *fast in practice*?

# OUR RESULT

## The Smoothed-Striding Algorithm

### Key Features:

- ▶ linear work and polylogarithmic span  
(like the Standard Algorithm)
- ▶ fast in practice  
(like the Strided Algorithm)
- ▶ theoretically optimal cache behavior  
(unlike any past algorithm)

# STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

## Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

- ▶ Good cache behavior in practice
- ▶ Worst case span is  $T_\infty \approx n$
- ▶ On random inputs span is  $T_\infty = \tilde{O}(n^{2/3})$

# STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

## Strided Algorithm


[Francis and Pannan, 92; Frias and Petit, 08]

- ▶ Good cache behavior in practice
- ▶ Worst case span is  $T_\infty \approx n$
- ▶ On random inputs span is  $T_\infty = \tilde{O}(n^{2/3})$

## Smoothed-Striding Algorithm


- ▶ Provably optimal cache behavior
- ▶ Span is  $T_\infty = O(\log n \log \log n)$  with high probability in  $n$
- ▶ Uses randomization *inside* the algorithm

# SMOOTHED-STRIDING ALGORITHM'S PERFORMANCE




# The Smoothed-Striding Algorithm

Logically partition the array into chunks of adjacent elements




Form groups  $U_i$  that contain a **random** element from each chunk (the Strided Algorithm forms groups  $P_i$  that contain the  $i$ -th element from each chunk)






Perform serial partitions on each  $U_i$  in parallel over the  $U_i$ 's




This step is highly parallel.


Define  $v_i = \text{index of first element greater than the pivot in } U_i$




Identify leftmost and rightmost  $v_i$



**Final step:** Recursively partition the subarray



**Final step:** Recursively partition the subarray




- ▶ Recursion is now possible!  
(unlike in the Strided Algorithm, where the subproblem is a worst case input)
- ▶ Randomness guarantees that  $v_{\max} - v_{\min}$  is small  
(unlike in the Strided Algorithm, where  $v_{\max} - v_{\min}$  could be as large as  $\Theta(n)$ )

## A KEY CHALLENGE


How do we store the  $U_i$ 's if they are all random?

Storing which elements make up each  $U_i$  takes too much space!

**Strided Algorithm  $P_i$ .**




**Smoothed-Striding Algorithm  $U_i$ .**



# HOW TO STORE THE GROUPS

**Key Insight:** While each  $U_i$  does need to contain a random element from each chunk, the  $U_i$ 's don't need to be *independent*.

We store  $U_1$ , and all other groups are determined by a “circular shift” of  $U_1$  (wraparound within each chunk).



## AN OPEN QUESTION

**Our algorithm:**  $\text{span } T_\infty = O(\log n \log \log n)$

**Standard Algorithm:**  $\text{span } T_\infty = O(\log n)$ .

Can we get optimal cache behavior and  $\text{span } O(\log n)$ ?