

# Cache-Efficient Parallel Partition Algorithms Using Exclusive-Read-and-Write Memory


William Kuszmaul, Alek Westover

MIT

July 4, 2020


# THE PARTITION PROBLEM

An unpartitioned array:




# THE PARTITION PROBLEM

An unpartitioned array:



An array partitioned relative to a pivot value:




## EXAMPLE APPLICATIONS

- ▶ Parallel Partition
- ▶ Parallel Quicksort
- ▶ Filtering Operations


# OUR MODEL OF PARALLELISM

Language based model. Primitive: *Parallel for loop*.  
Similar to how Cilk works



# OUR MODEL OF PARALLELISM

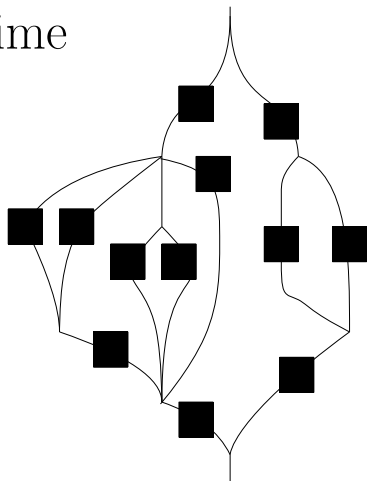
Language based model. Primitive: *Parallel for loop*.  
Similar to how Cilk works



Rule for this talk: no locks or atomic variables

# PARALLEL ALGORITHM PERFORMANCE METRICS

Time



$T_p$ : Time to run on  $p$  processors

**Work:**  $T_1$

► time to run in serial

**Span:**  $T_\infty$

► time to run on infinitely many processors

Brent's Theorem:

$$T_p = \Theta \left( \frac{T_1}{p} + T_\infty \right)$$

# OPTIMIZING CACHE EFFICIENCY

*Cache-efficient* algorithm: serial implementation incurs  $(1 + o(1))n/B$  cache misses in the Disk Access model.

[Aggarwal and Vitter, 1988]

- ▶ Perform low number of passes over the data
- ▶ Don't use extra memory (*In-Place*)
- ▶ Simultaneously operate on elements close in memory



## PREVIOUS WORK

“Standard Algorithm”: span  $O(\log n)$ , but slow in practice

- ▶ Uses extra memory
  - ▶ Makes multiple passes over data
- } bad cache behavior

Fastest algorithms in practice:

- ▶ Lock based and atomic-variable based algorithms

[Michael Axtmann, Sascha Witt, Daniel Ferizovic, and Peter Sanders, 2017; Philip Heidelberger, Alan Norton, and John T. Robinson, 1990; Philippas Tsigas and Yi Zhang, 2003]

- ▶ The Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

No locks or atomic-variables, lacks theoretical guarantees

## OUR QUESTION

Can we create an algorithm with *theoretical guarantees* that is *fast in practice*?

# OUR RESULT

## The Smoothed-Striding Algorithm

### Key Features:

- ▶ linear work and polylogarithmic span  
(like the Standard Algorithm)
- ▶ fast in practice  
(like the Strided Algorithm)
- ▶ theoretically optimal cache behavior  
(unlike any past algorithm)

# STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

## Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

- ▶ Good cache behavior in practice
- ▶ Worst case span is  $T_\infty \approx n$
- ▶ On random inputs span is  $T_\infty = \tilde{O}(n^{2/3})$

# STRIDED VERSUS SMOOTHED-STRIDING ALGORITHM

## Strided Algorithm

[Francis and Pannan, 92; Frias and Petit, 08]

- ▶ Good cache behavior in practice
- ▶ Worst case span is  $T_\infty \approx n$
- ▶ On random inputs span is  $T_\infty = \tilde{O}(n^{2/3})$

## Smoothed-Striding Algorithm

- ▶ Provably optimal cache behavior
- ▶ Span is  $T_\infty = O(\log n \log \log n)$  with high probability in  $n$
- ▶ Uses randomization *inside* the algorithm

# SMOOTHED-STRIDING ALGORITHM'S PERFORMANCE

